

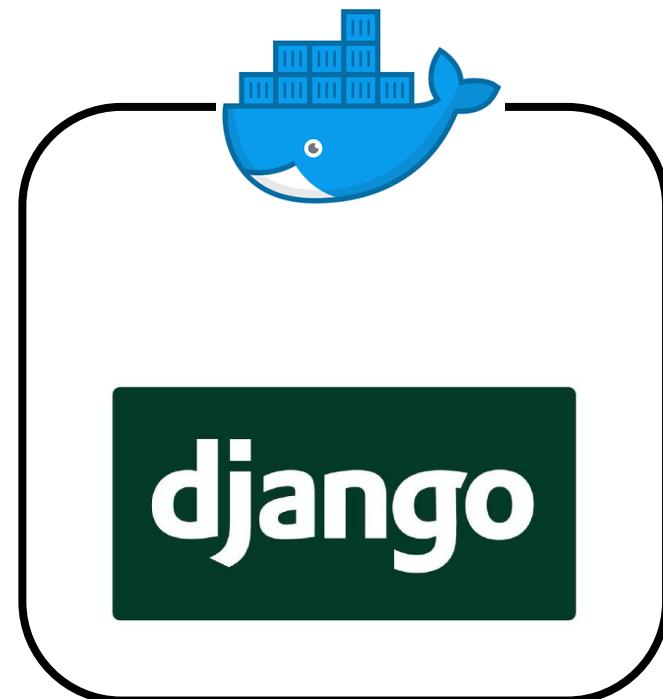
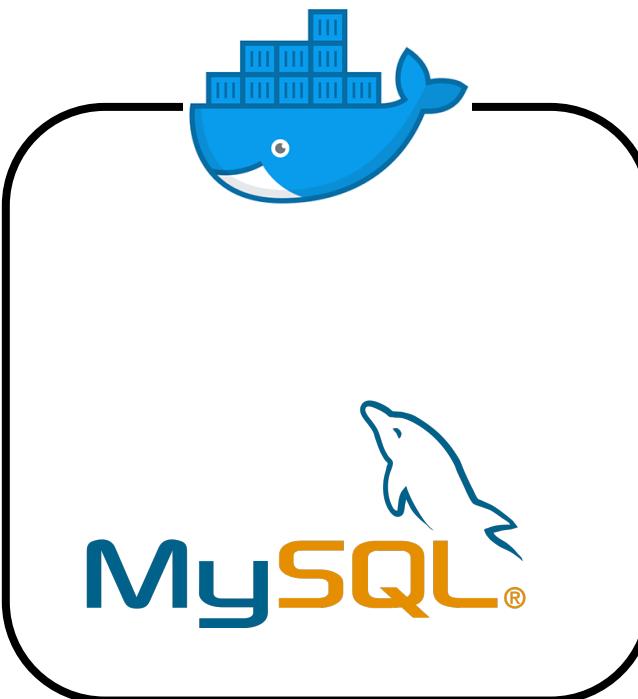
MLOps를 위한 Kubernetes 기초



kubernetes

하나의 서비스를 만드는데 필요한 컨테이너는?

간단한 웹 서비스를 만든다면?



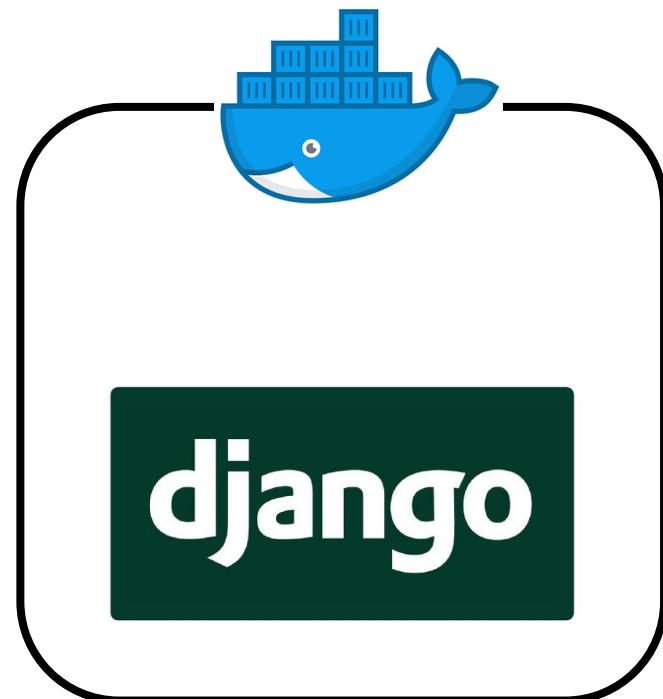
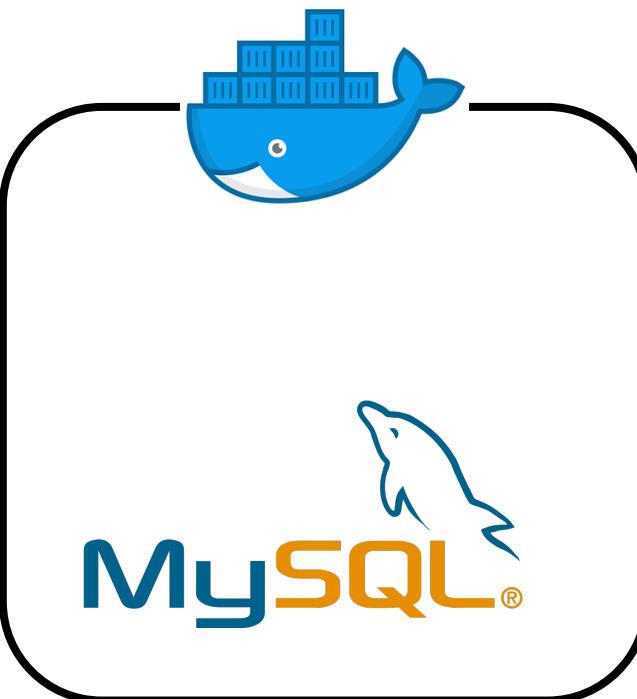
```
docker run --name mysql -p 3306:3306 mysql:latest
```

```
docker run --name myapp -p 8080:80 myapp:django
```

```
docker run --name nginx -p 80:80 nginx:latest
```

하나의 서비스를 만드는데 필요한 컨테이너는?

간단한 웹 서비스를 만든다면?



```
docker run --name mysql -p 3306:3306 mysql:latest
```

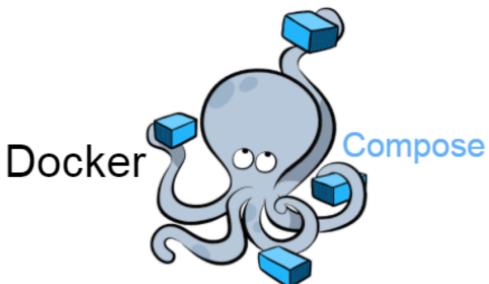
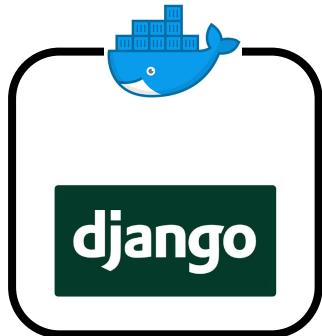
```
docker run --name myapp -p 8080:80 myapp:django
```

```
docker run --name nginx -p 80:80 nginx:latest
```

뭔가 귀찮다
순서도 지켜야 한다.

하나의 서비스를 만드는데 필요한 컨테이너는?

간단한 웹 서비스를 만든다면?

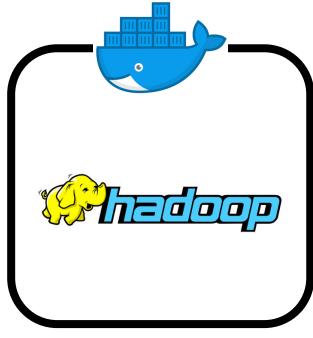
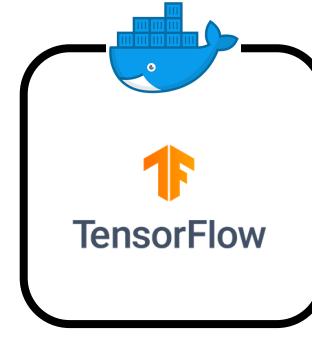
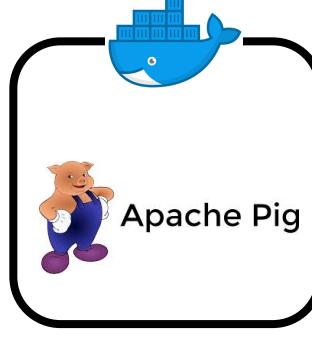
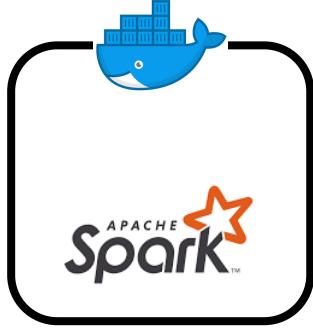
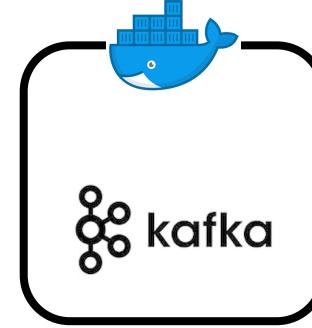
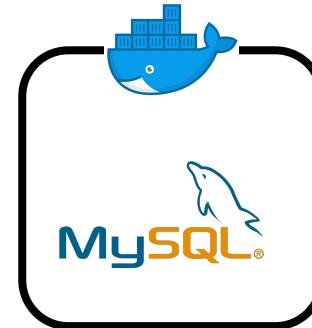


docker-compose up

```
python:  
  build: ./python  
  command: uwsgi --socket :8001 --module app.wsgi --py-autoreload 1 --logto /tmp/mylog.log  
  volumes:  
    - ./src:/code  
    - ./static:/static  
  expose:  
    - "8001"  
  depends_on:  
    - db  
  
db:  
  image: mysql:5.7  
  command: mysqld --character-set-server=utf8mb4 --collation-server=utf8mb4_unicode_ci  
  ports:  
    - "3306:3306"  
  environment:  
    MYSQL_ROOT_PASSWORD: root  
    MYSQL_DATABASE: todoList  
    MYSQL_USER: user  
    MYSQL_PASSWORD: password  
    TZ: 'Asia/Tokyo'  
  volumes:  
    - ./mysql:/var/lib/mysql  
    - ./sql:/docker-entrypoint-initdb.d  
  
nginx:  
  image: nginx:1.13  
  ports:  
    - "8000:8000"  
  volumes:  
    - ./nginx/conf:/etc/nginx/conf.d  
    - ./nginx/uwsgi_params:/etc/nginx/uwsgi_params  
    - ./static:/static  
  depends_on:  
    - python
```

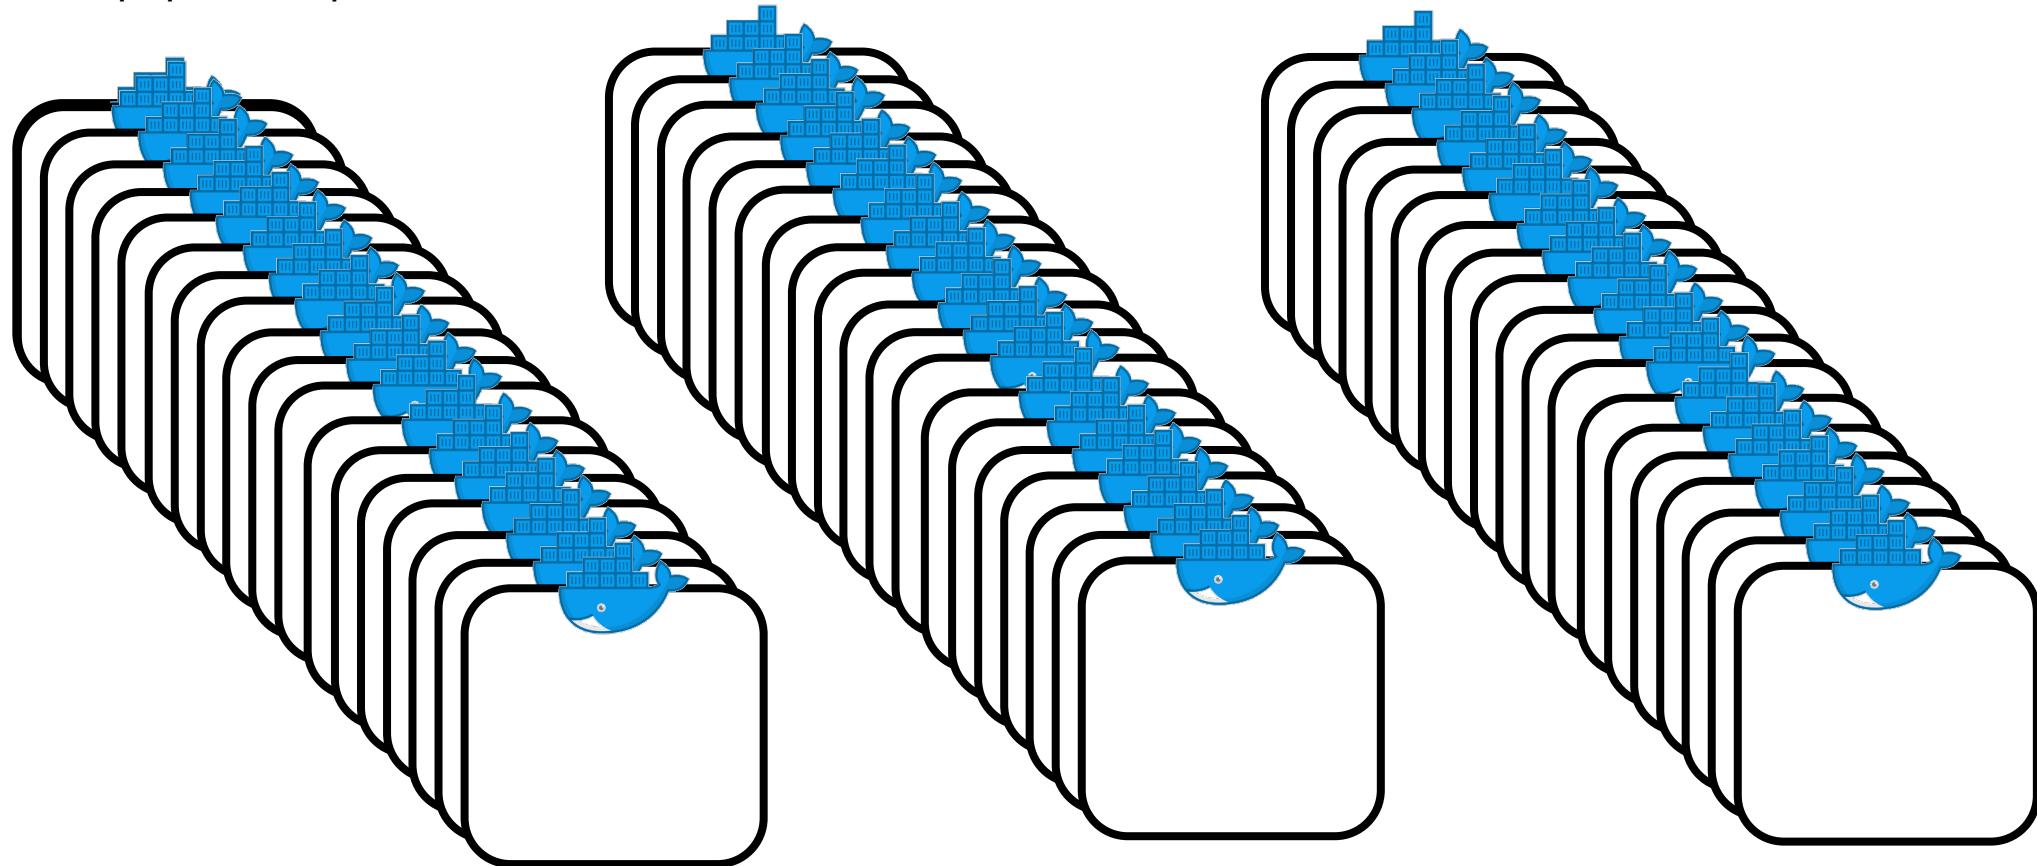
하나의 서비스를 만드는데 필요한 컨테이너는?

좀 더 복잡한 시스템이라면?



하나의 서비스를 만드는데 필요한 컨테이너는?

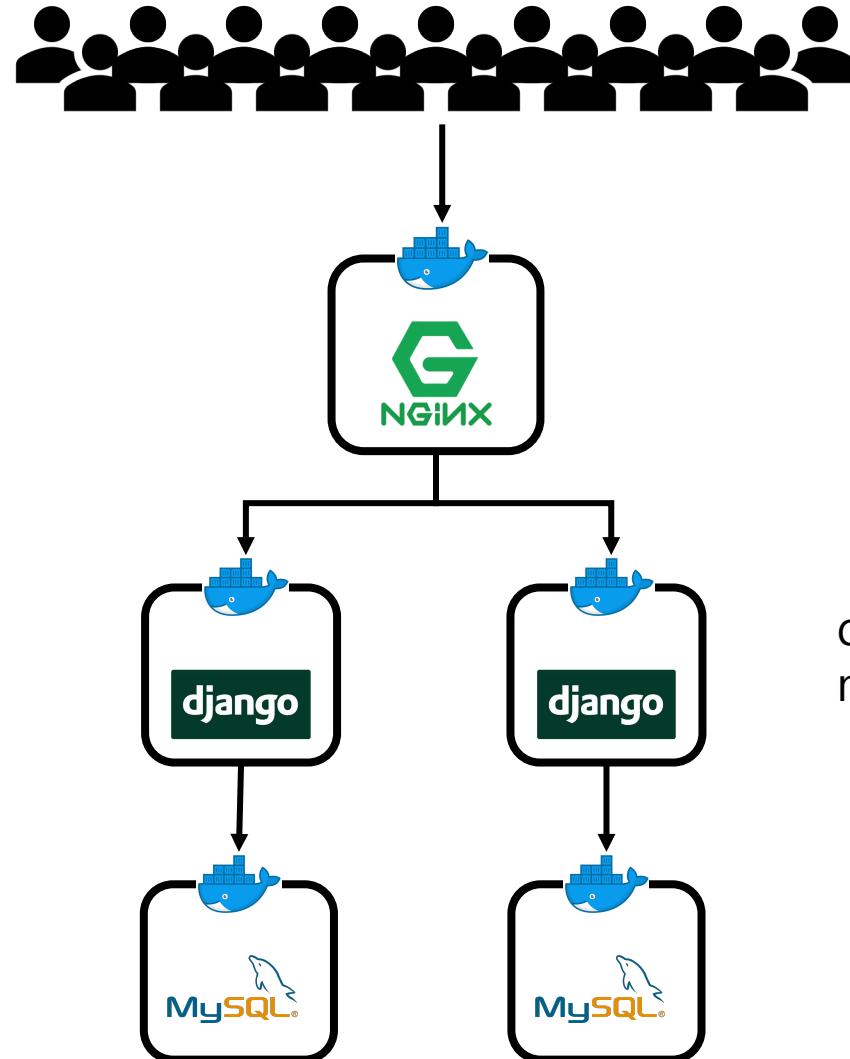
심지어 MSA라면?



Google은 1주에 약 1억 개 정도의 컨테이너를 띄운다고 합니다.

컨테이너의 관리

부하가 많이 생기면?

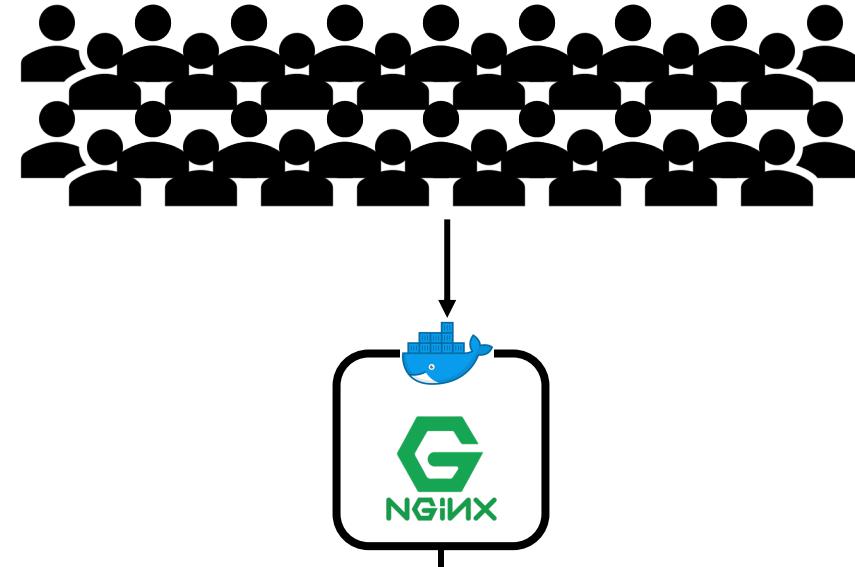


developer 또는 operation team

`docker-compose up --scale
myserver=2,db=2`

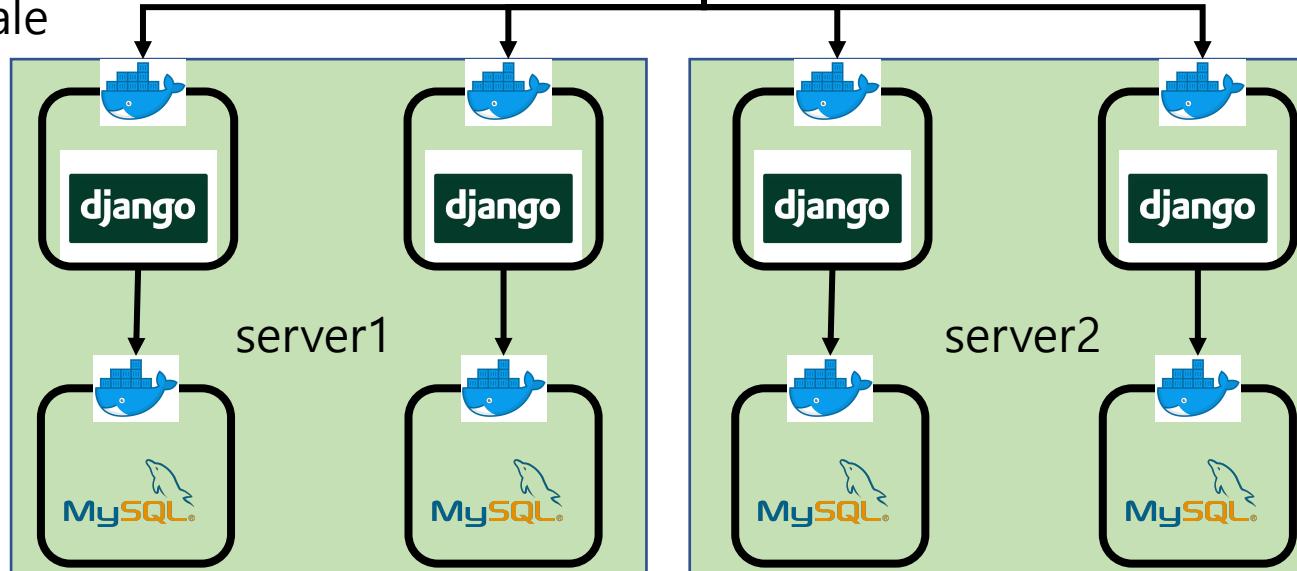
컨테이너의 관리

부하가 많이 생기면?



developer 또는 operation team

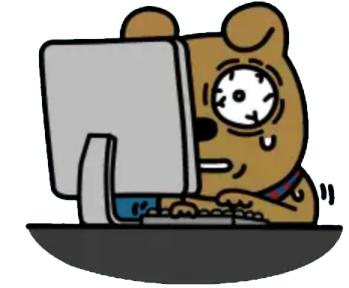
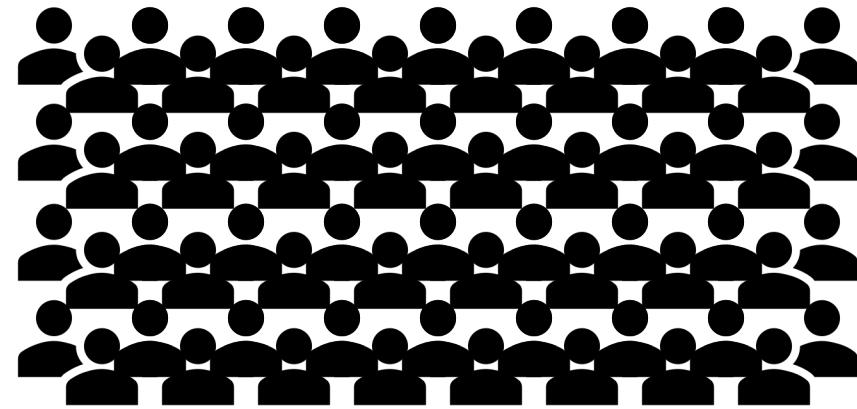
`docker-compose up --scale
myserver=2,db=2`



`docker-compose up --scale
myserver=2,db=2`

컨테이너의 관리

부하가 많이 생기면?



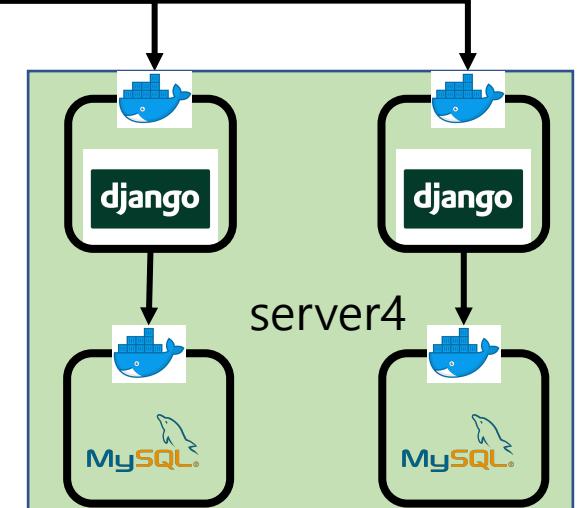
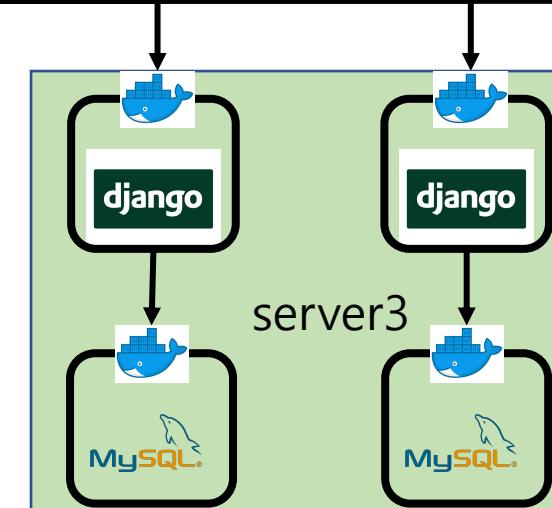
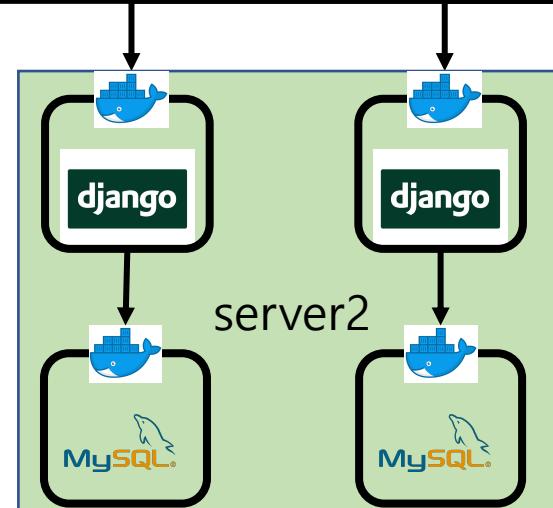
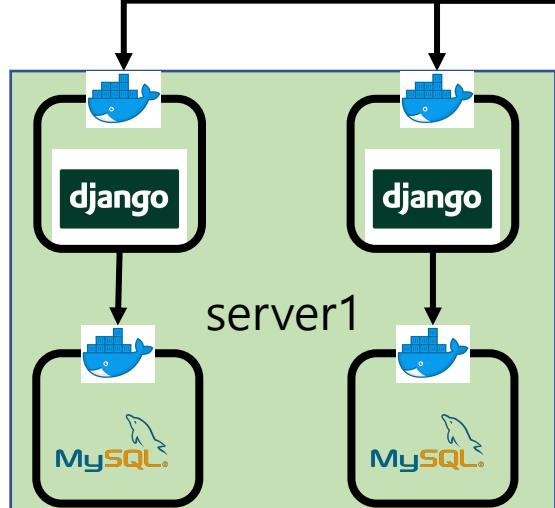
developer 또는 operation team

docker-compose up --scale
myserver=2,db=2

docker-compose up --scale
myserver=2,db=2

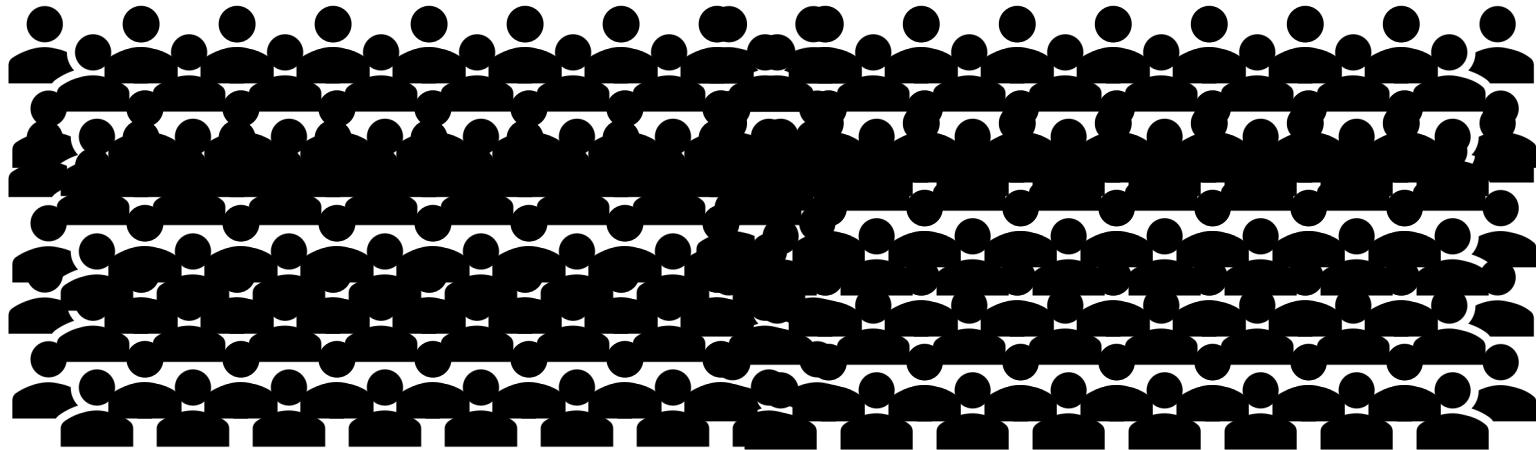
docker-compose up --scale
myserver=2,db=2

docker-compose up --scale
myserver=2,db=2



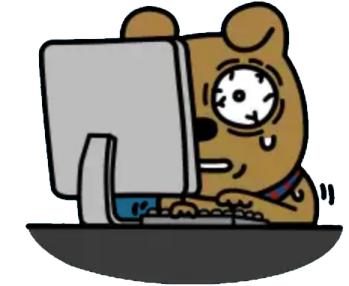
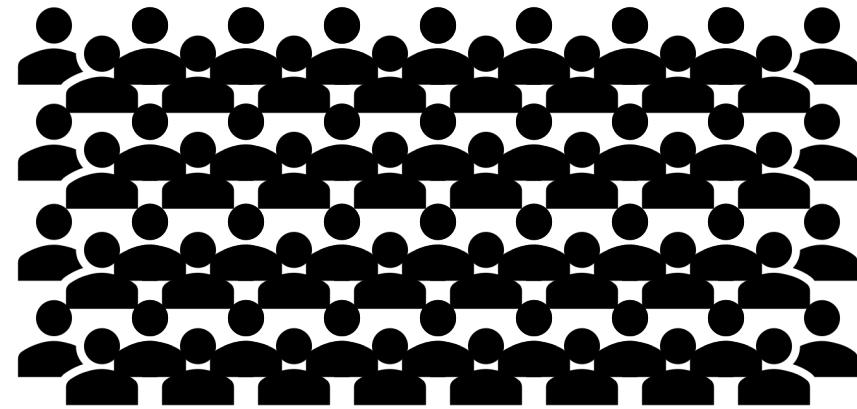
컨테이너의 관리

더 많아지면?

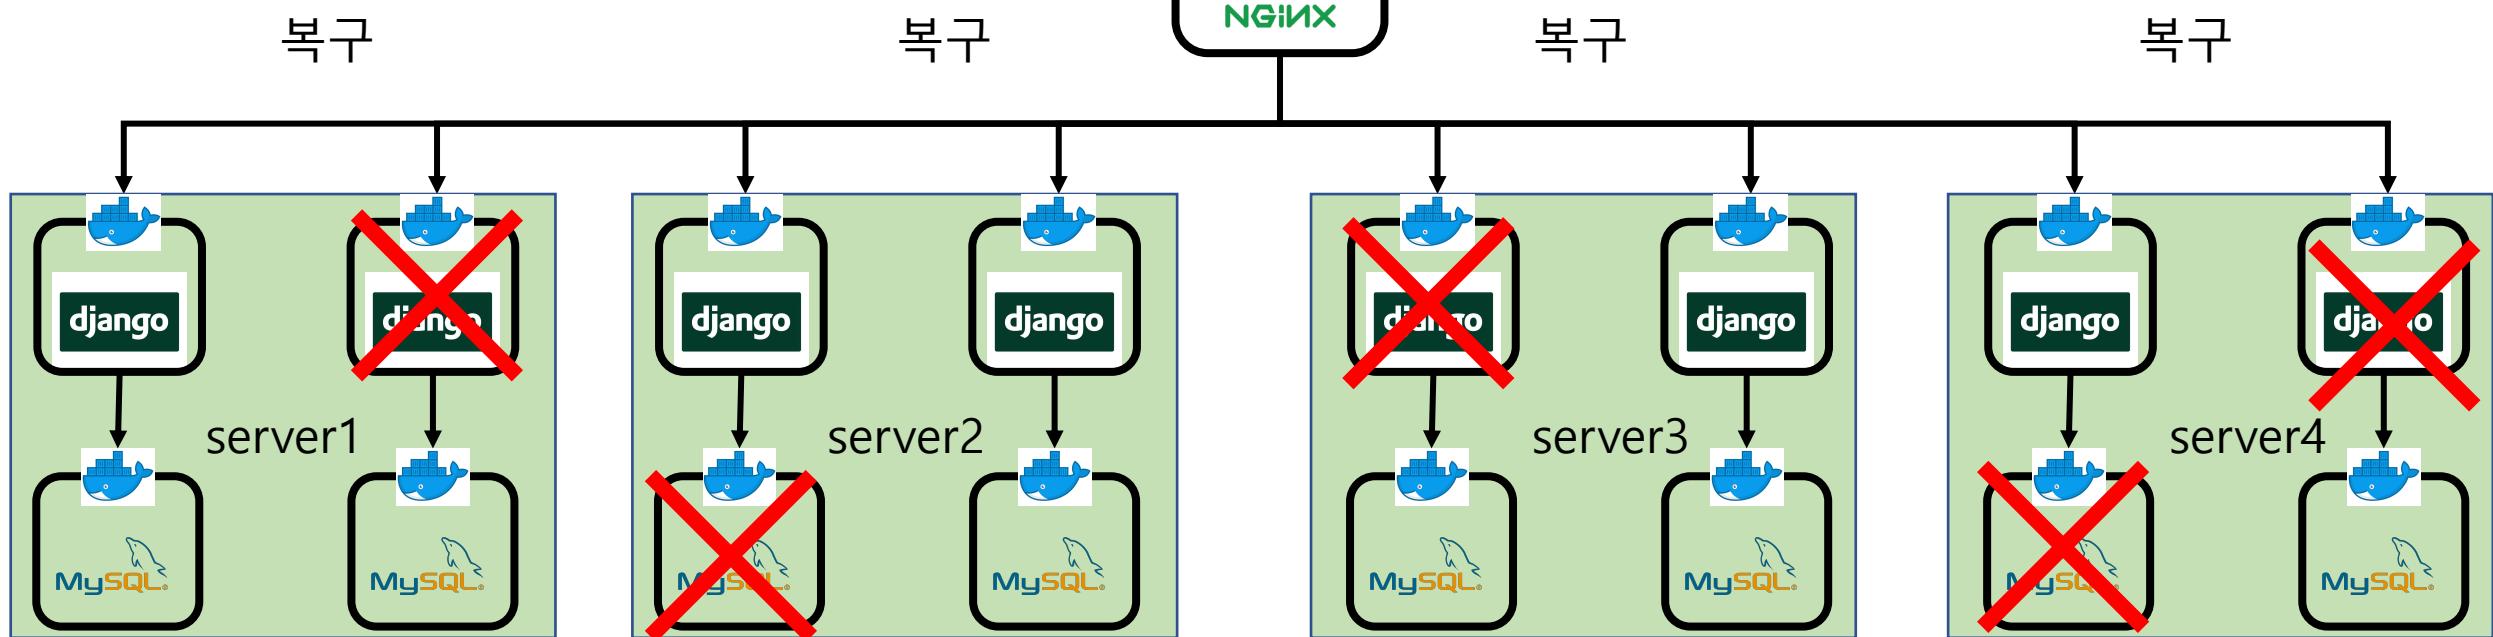


컨테이너의 관리

컨테이너가 망가지면?



developer 또는 operation team



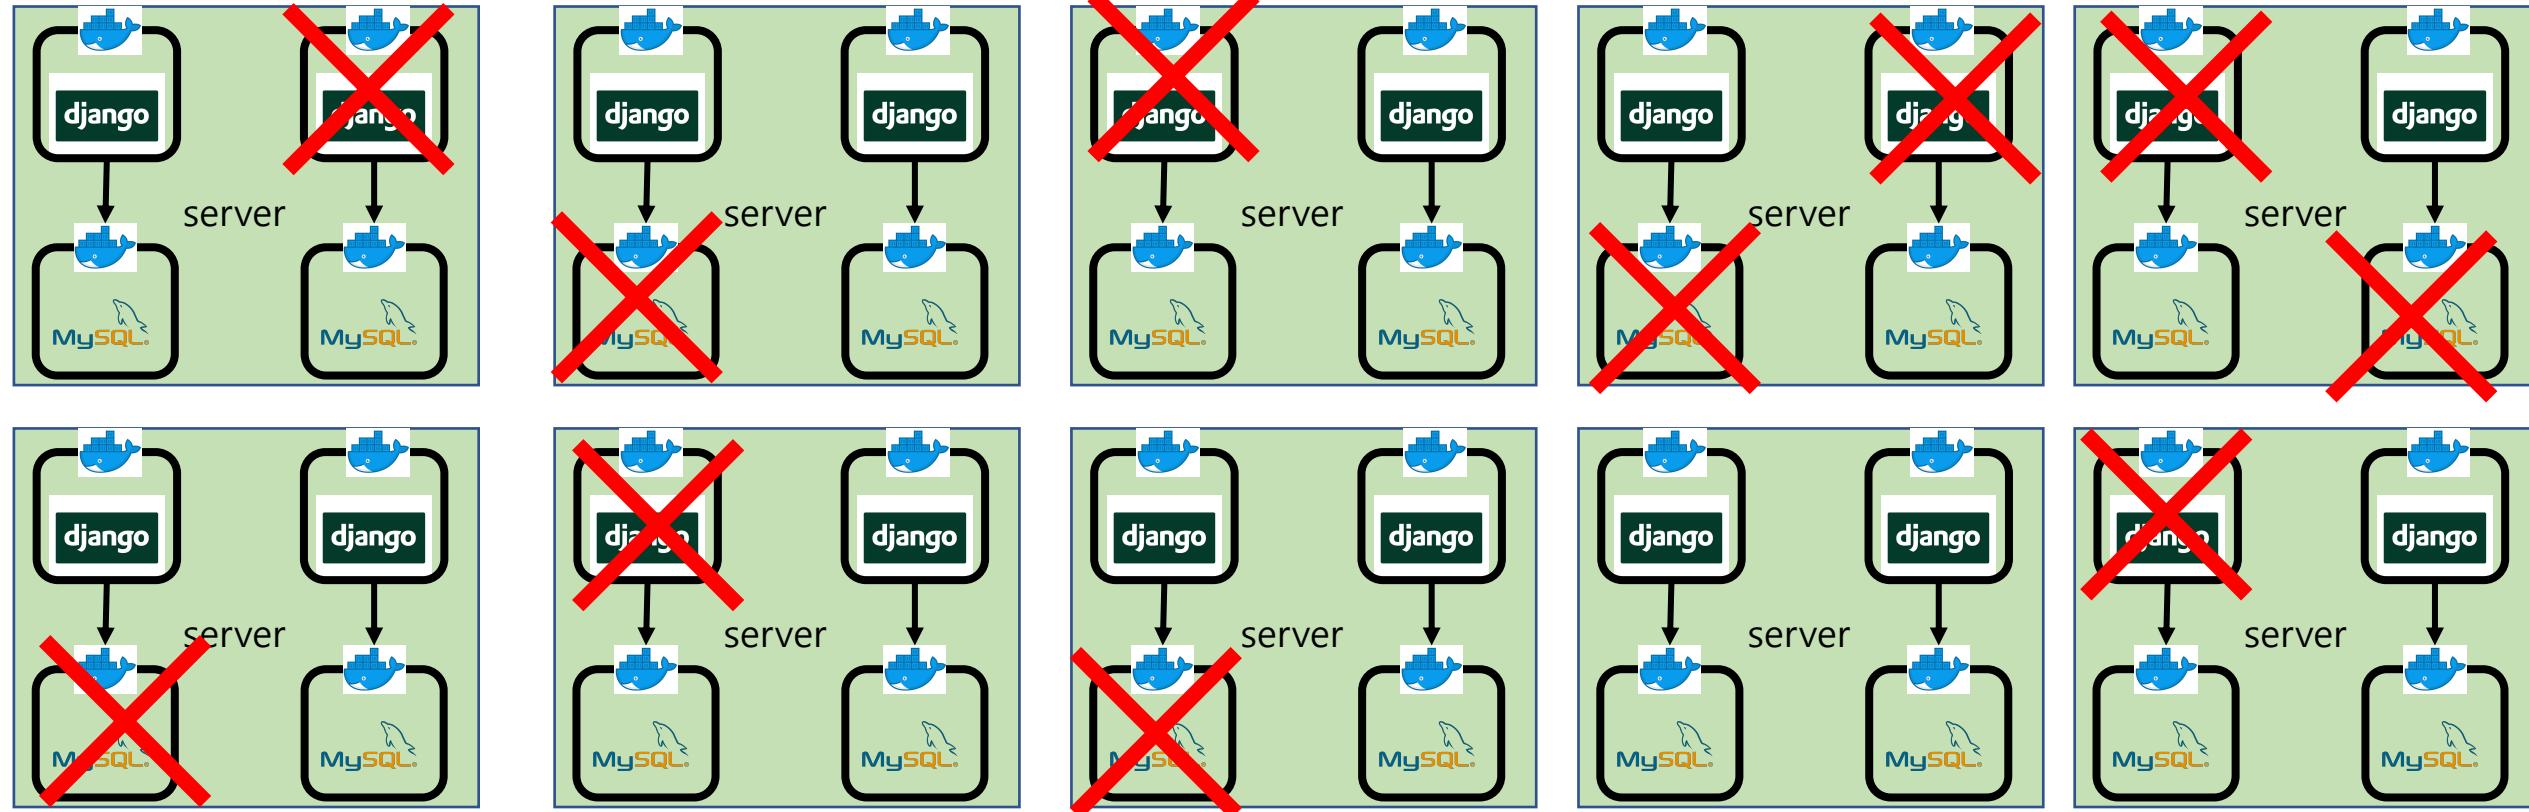
컨테이너의 관리

이거 일일히 다 복구하라구요?



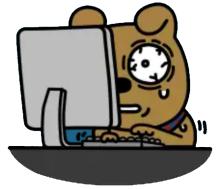
컨테이너가 더 많으면?

developer 또는 operation team

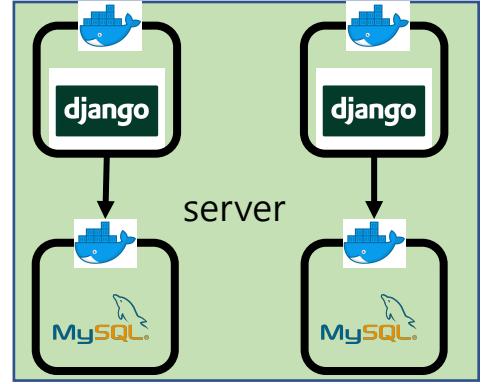


컨테이너의 관리

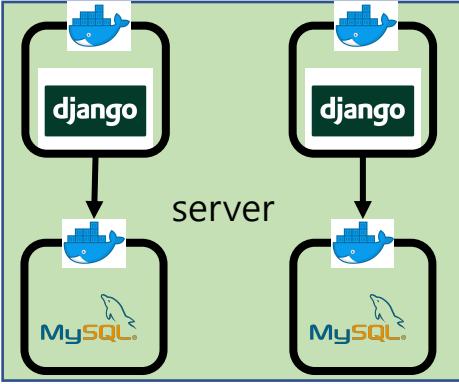
서버가 업데이트되면?



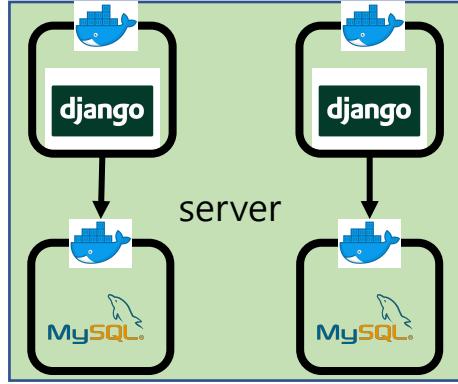
developer 또는 operation team



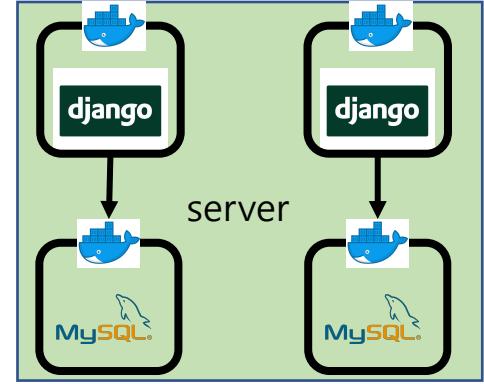
내리고 -> 이미지 pull -> 올리고



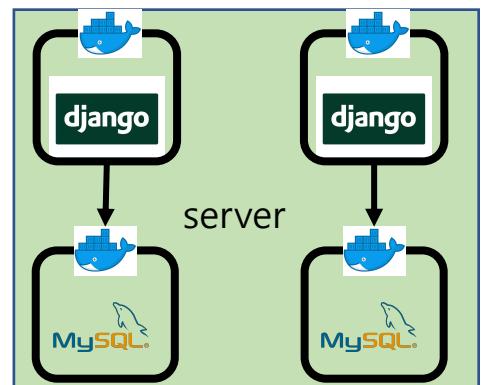
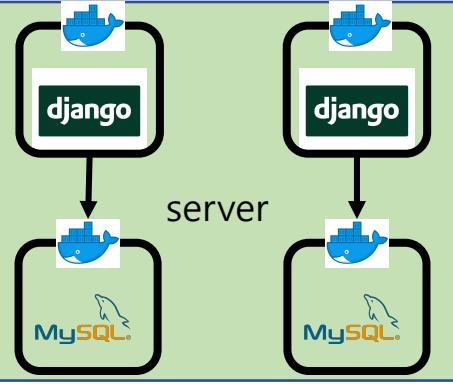
내리고 -> 이미지 pull -> 올리고



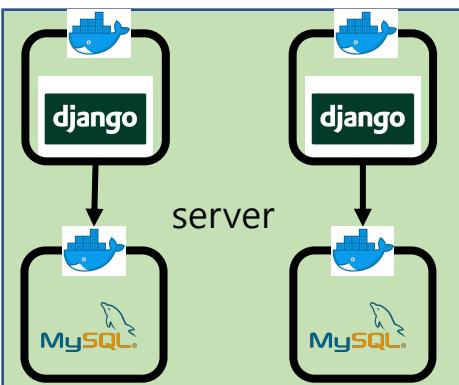
내리고 -> 이미지 pull -> 올리고



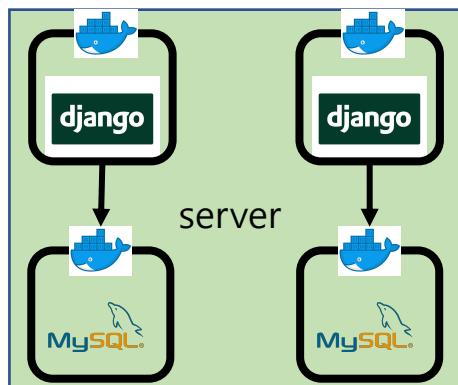
내리고 -> 이미지 pull -> 올리고 내리고 -> 이미지 pull -> 올리고



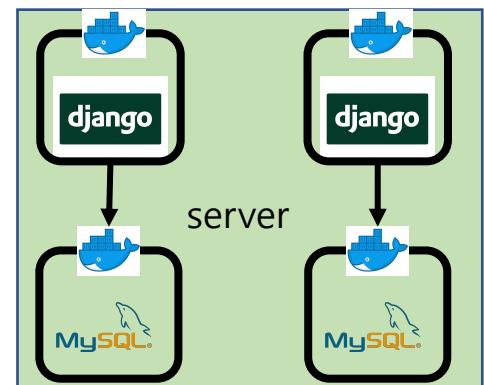
나는 누구



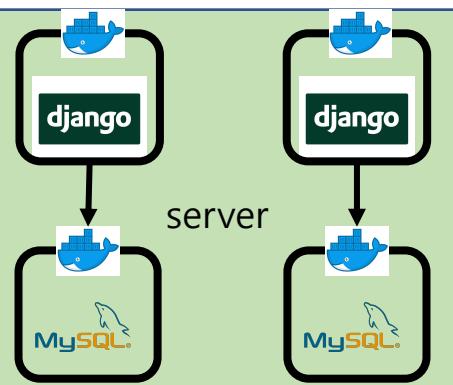
여긴 어디



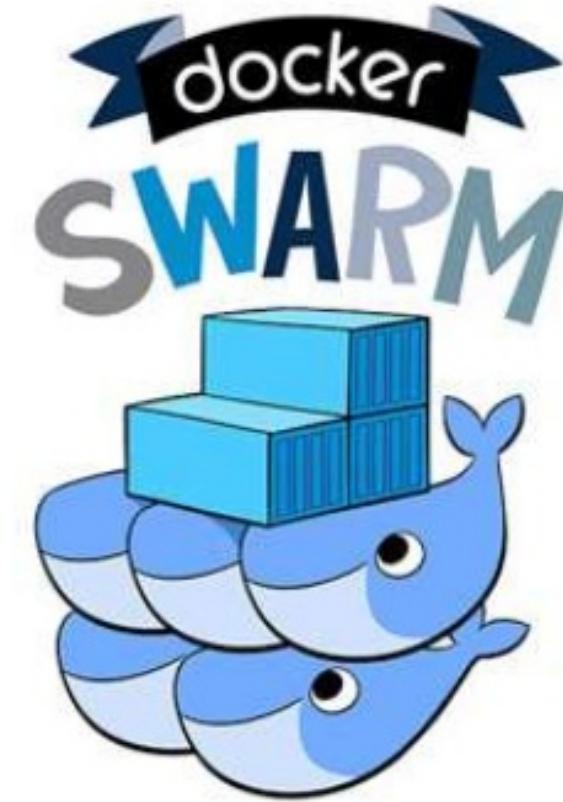
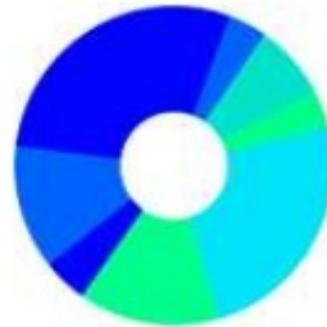
퇴사만이

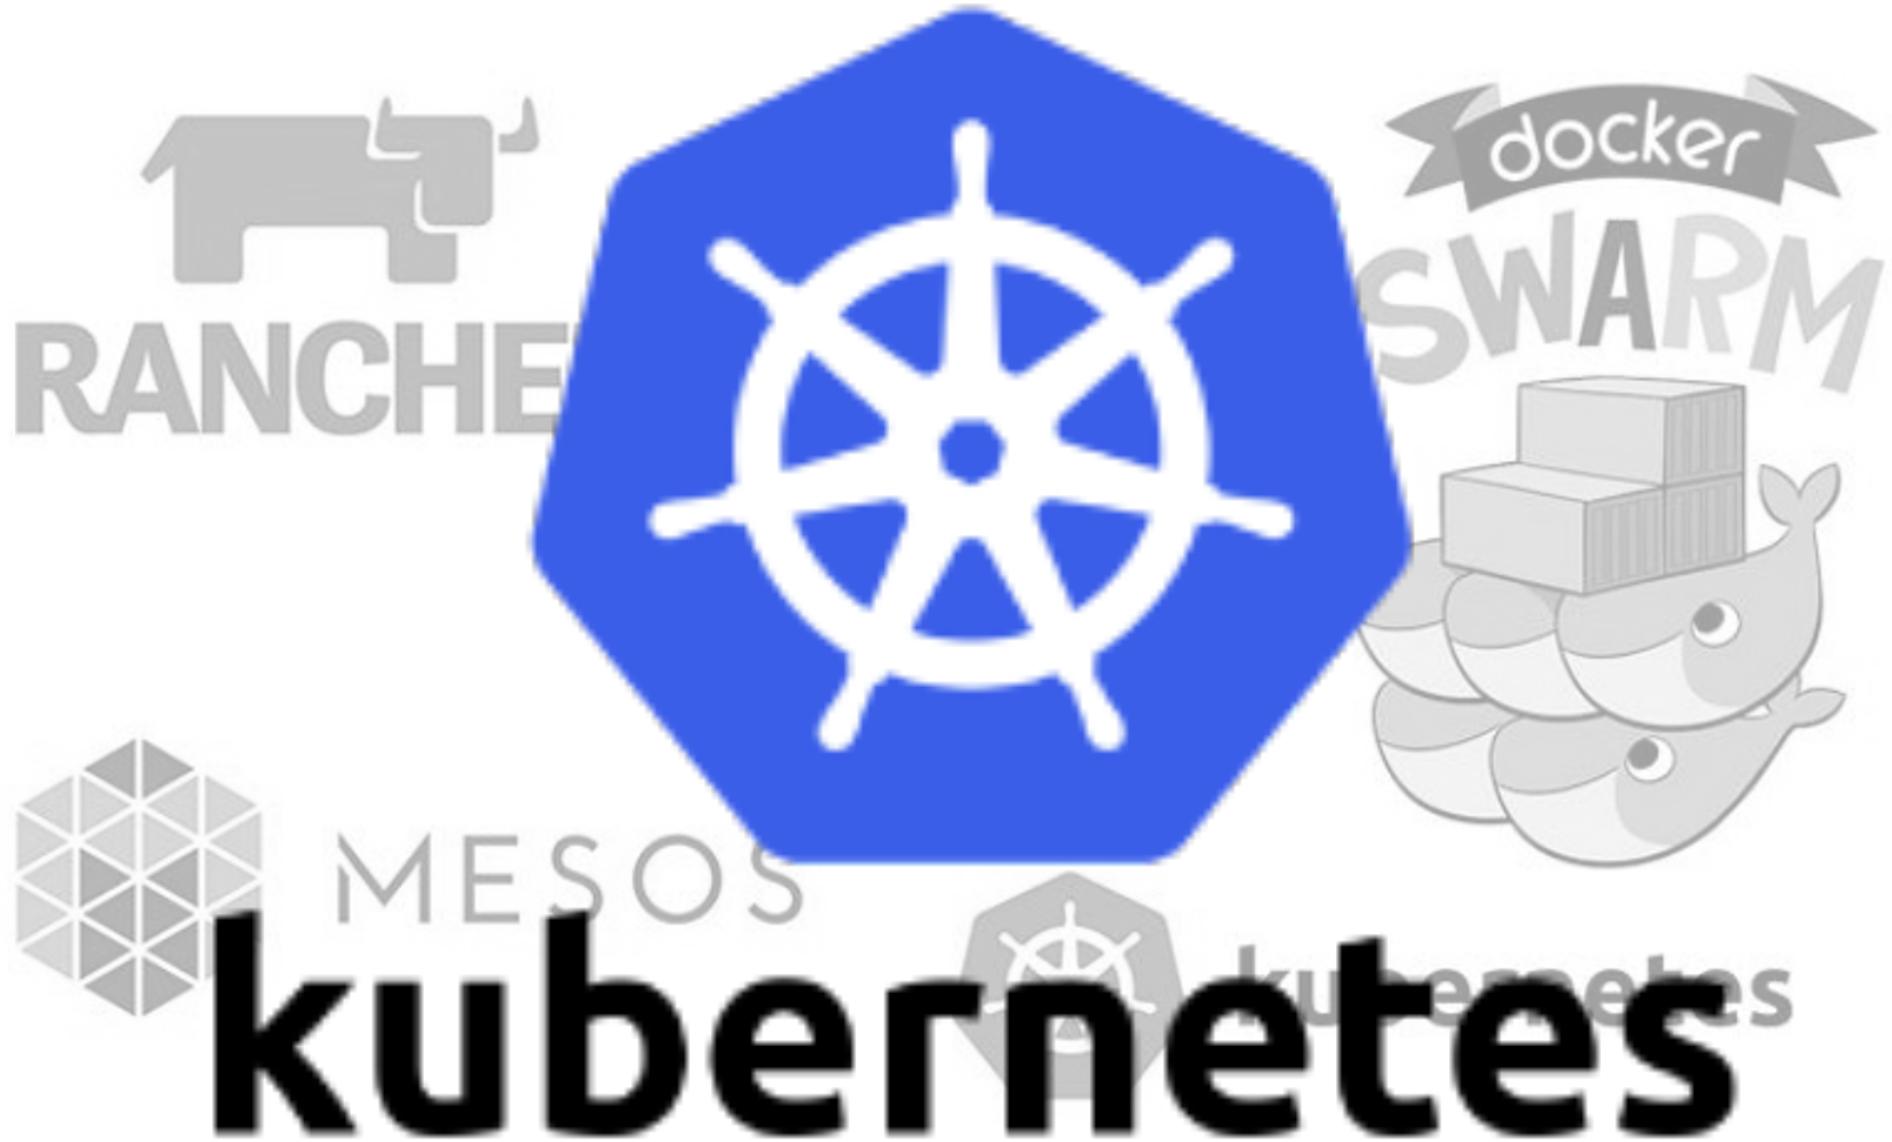


답인가



아아흑





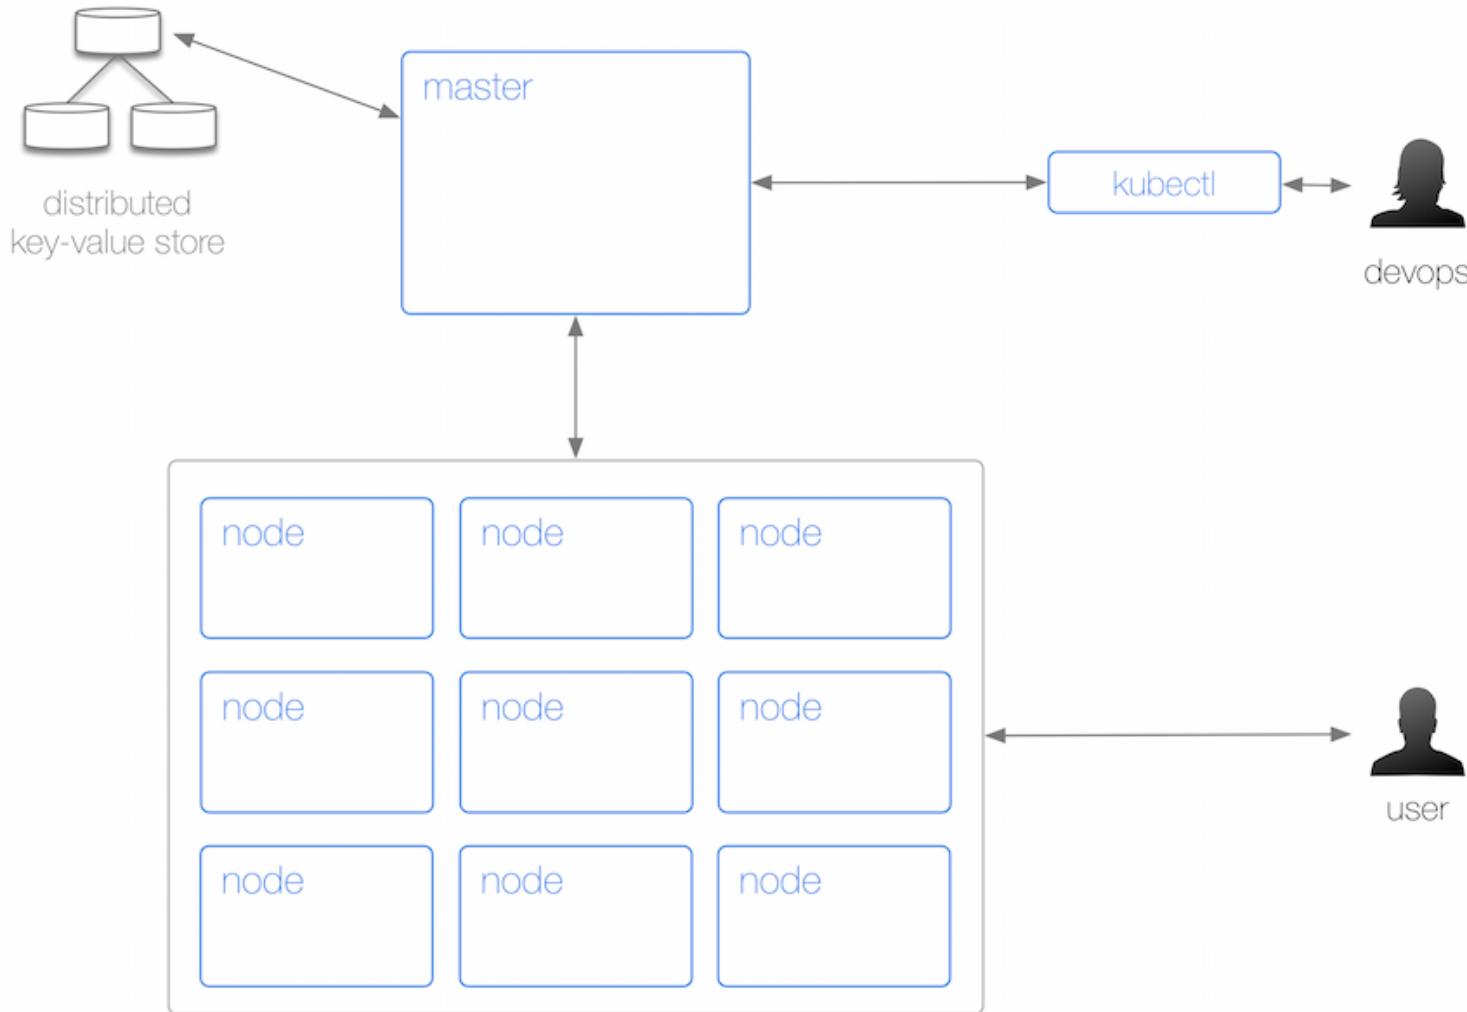
What does Orchestration do

- 여러 Host (Node)를 묶어 클러스터를 구성
- Container 를 적절한 위치에 배포 (Auto-placement)
- Container 가 죽으면 자동으로 복구 (Auto-restart)
- 필요에 따라 container를 Scaling, Replication, rolling update, rollback
- 기타 기능들([What is k8s?](#))

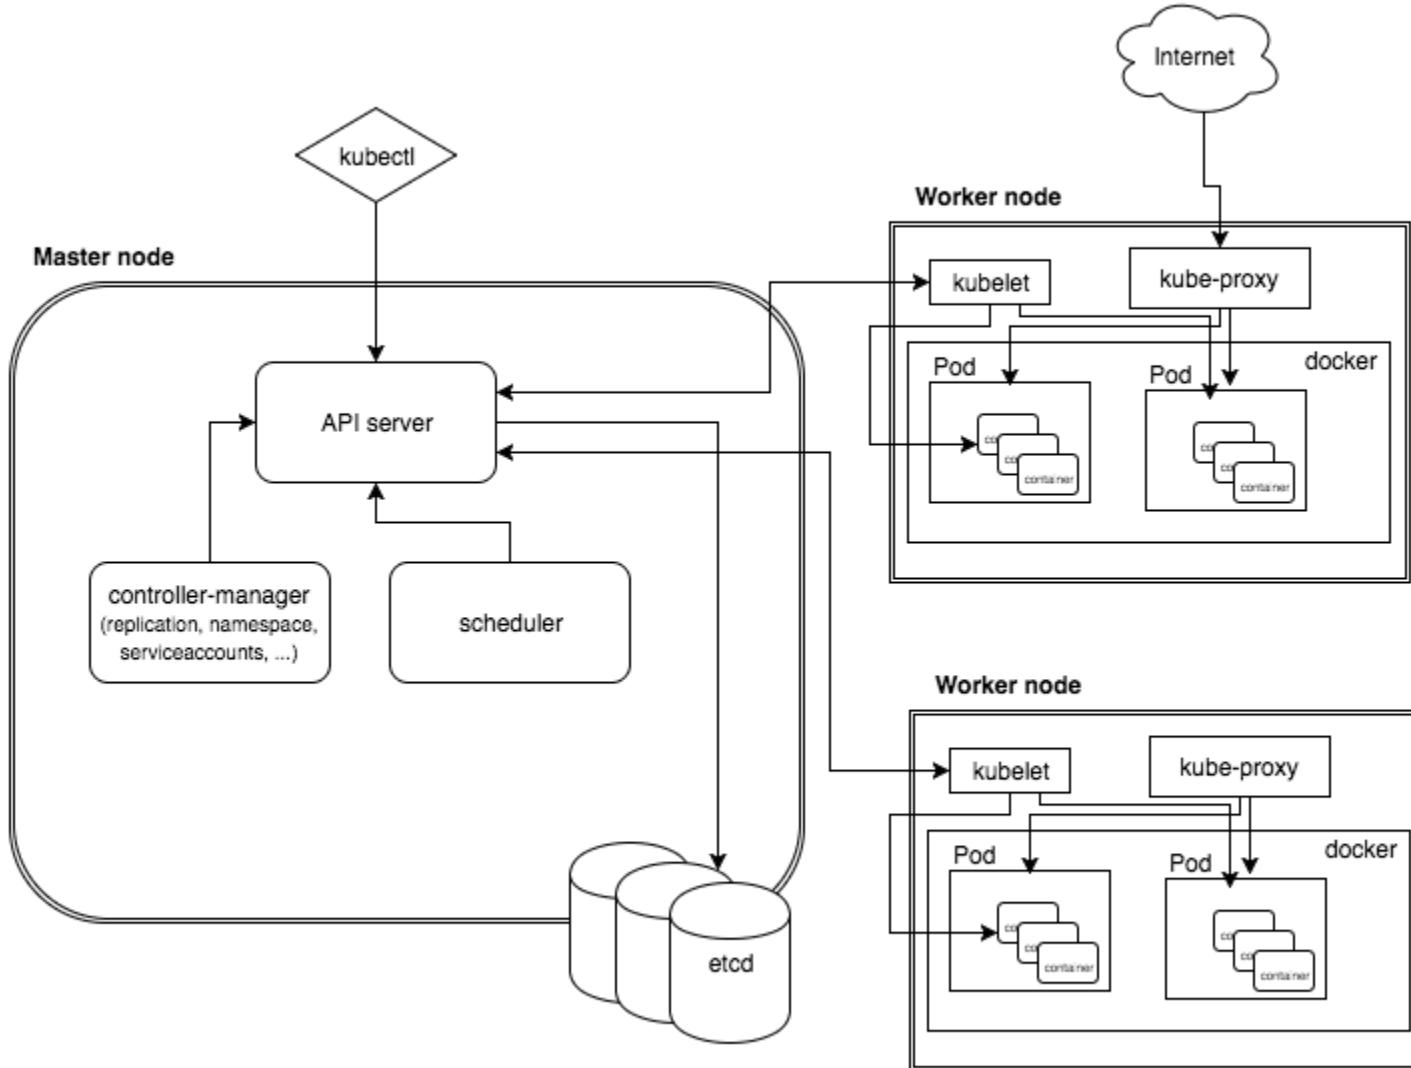
Should be understood to use

- kubernetes object
(e. g. pod, volume, service, controller)
- kubernetes architecture
- multi-host 위에서 container 실행 시 고려해야 할 것들
 - service discovery
 - networking
 - volume management
 - log aggregation

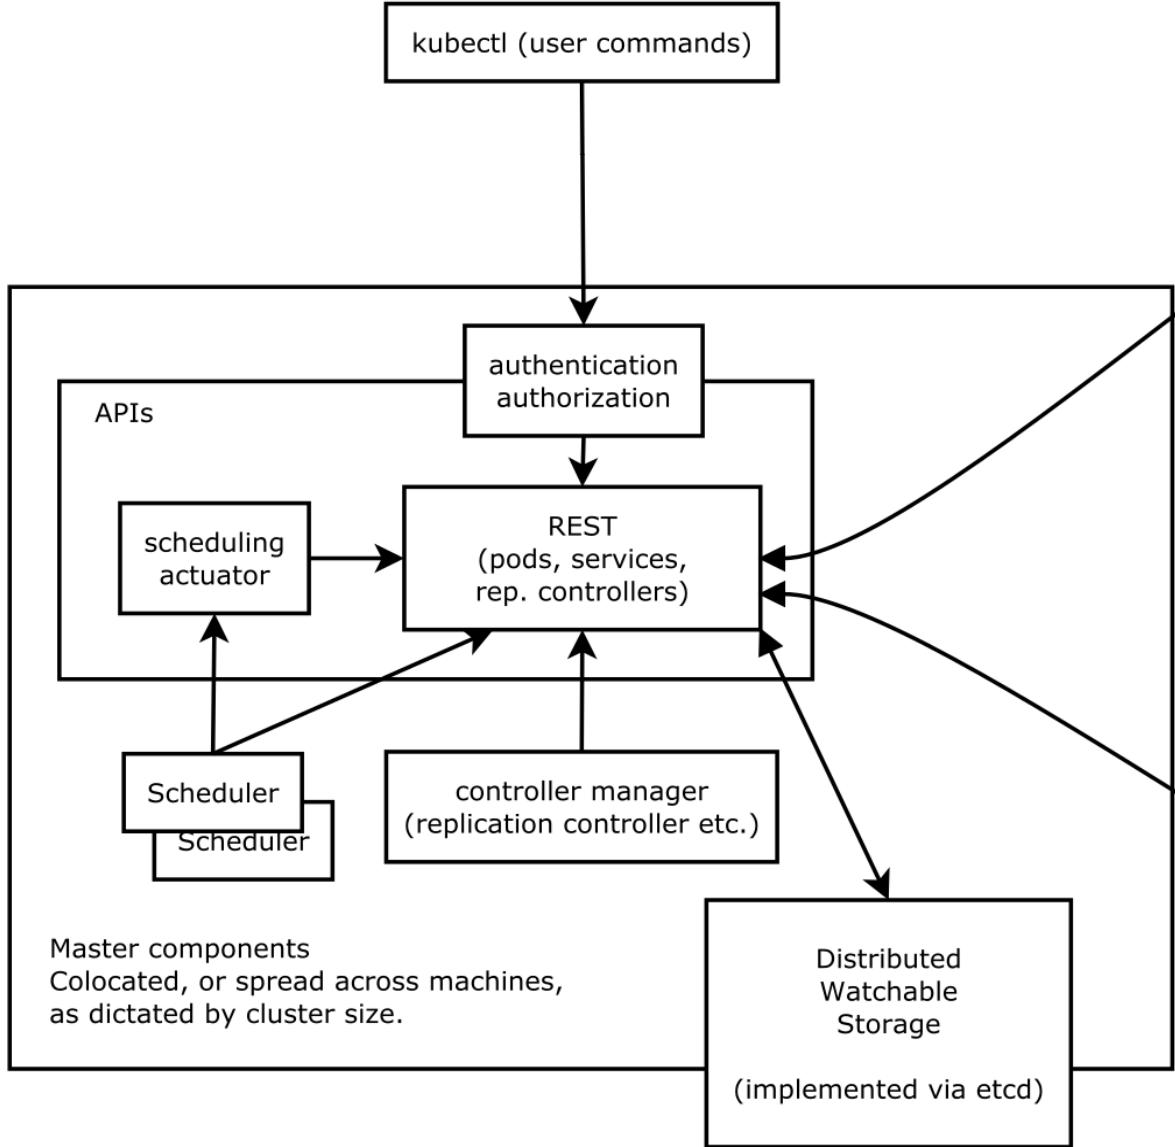
Kubernetes Architecture (abbreviated)



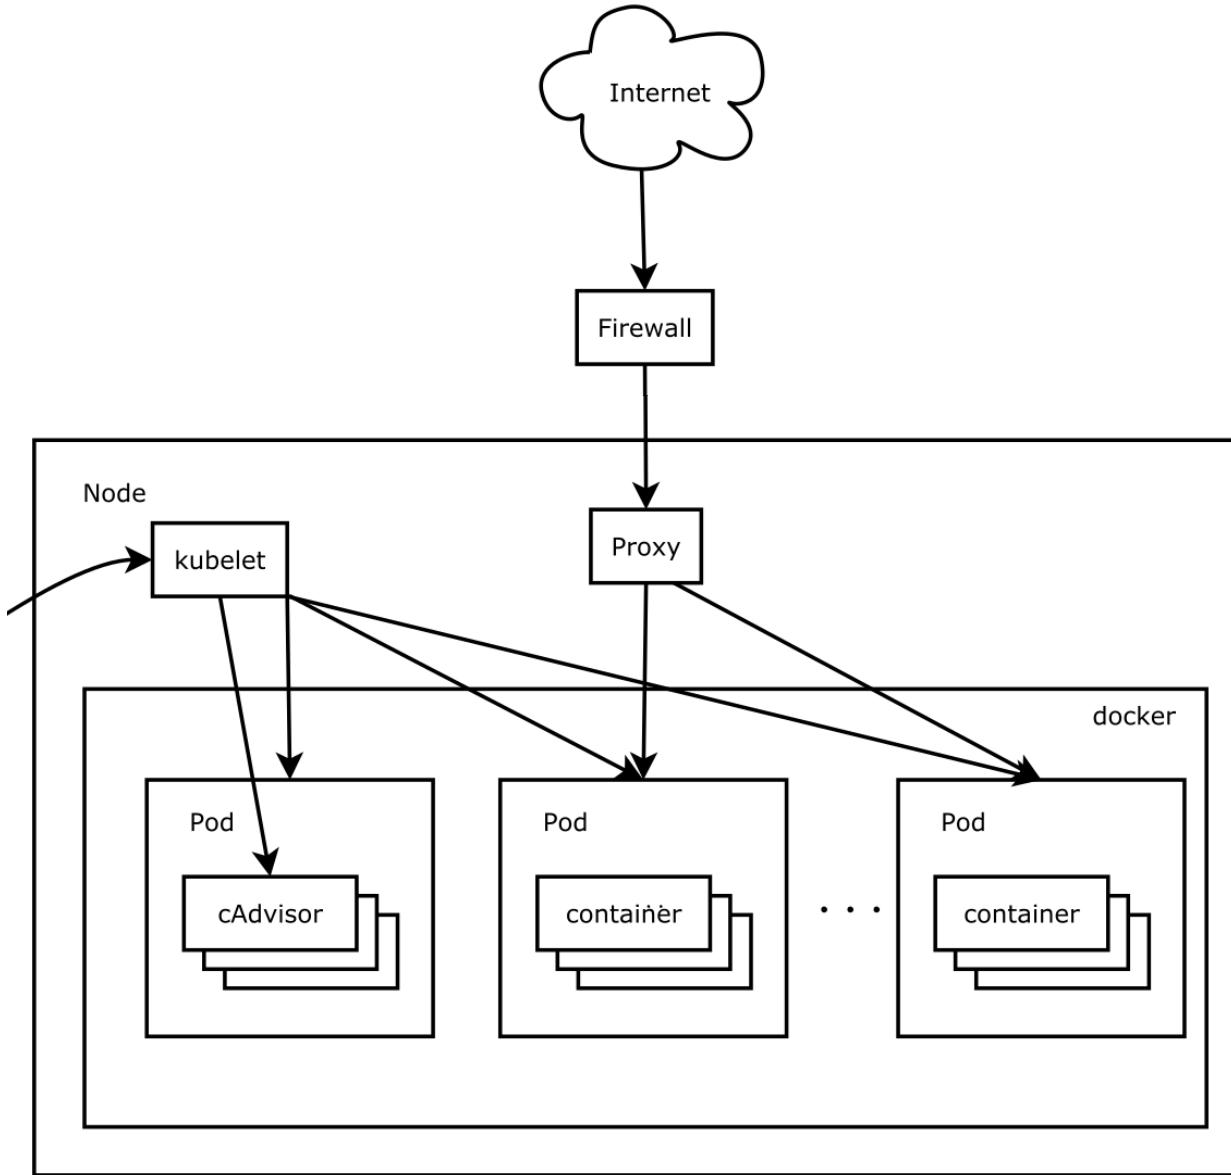
Kubernetes Architecture (detail)



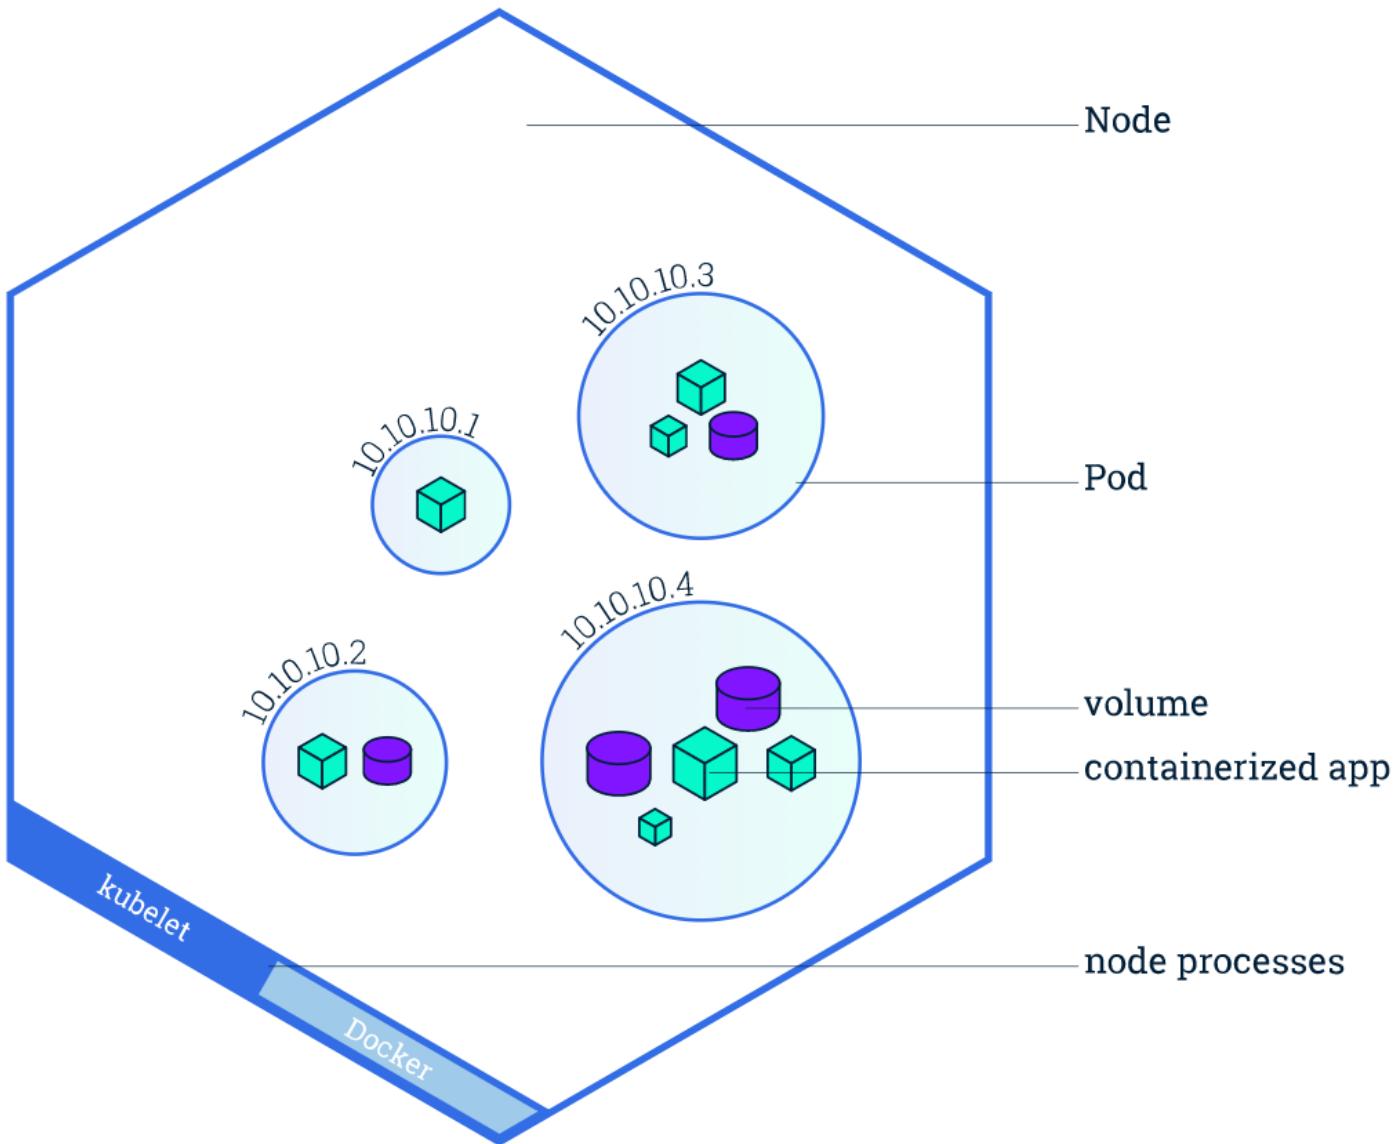
Kubernetes Architecture (master)



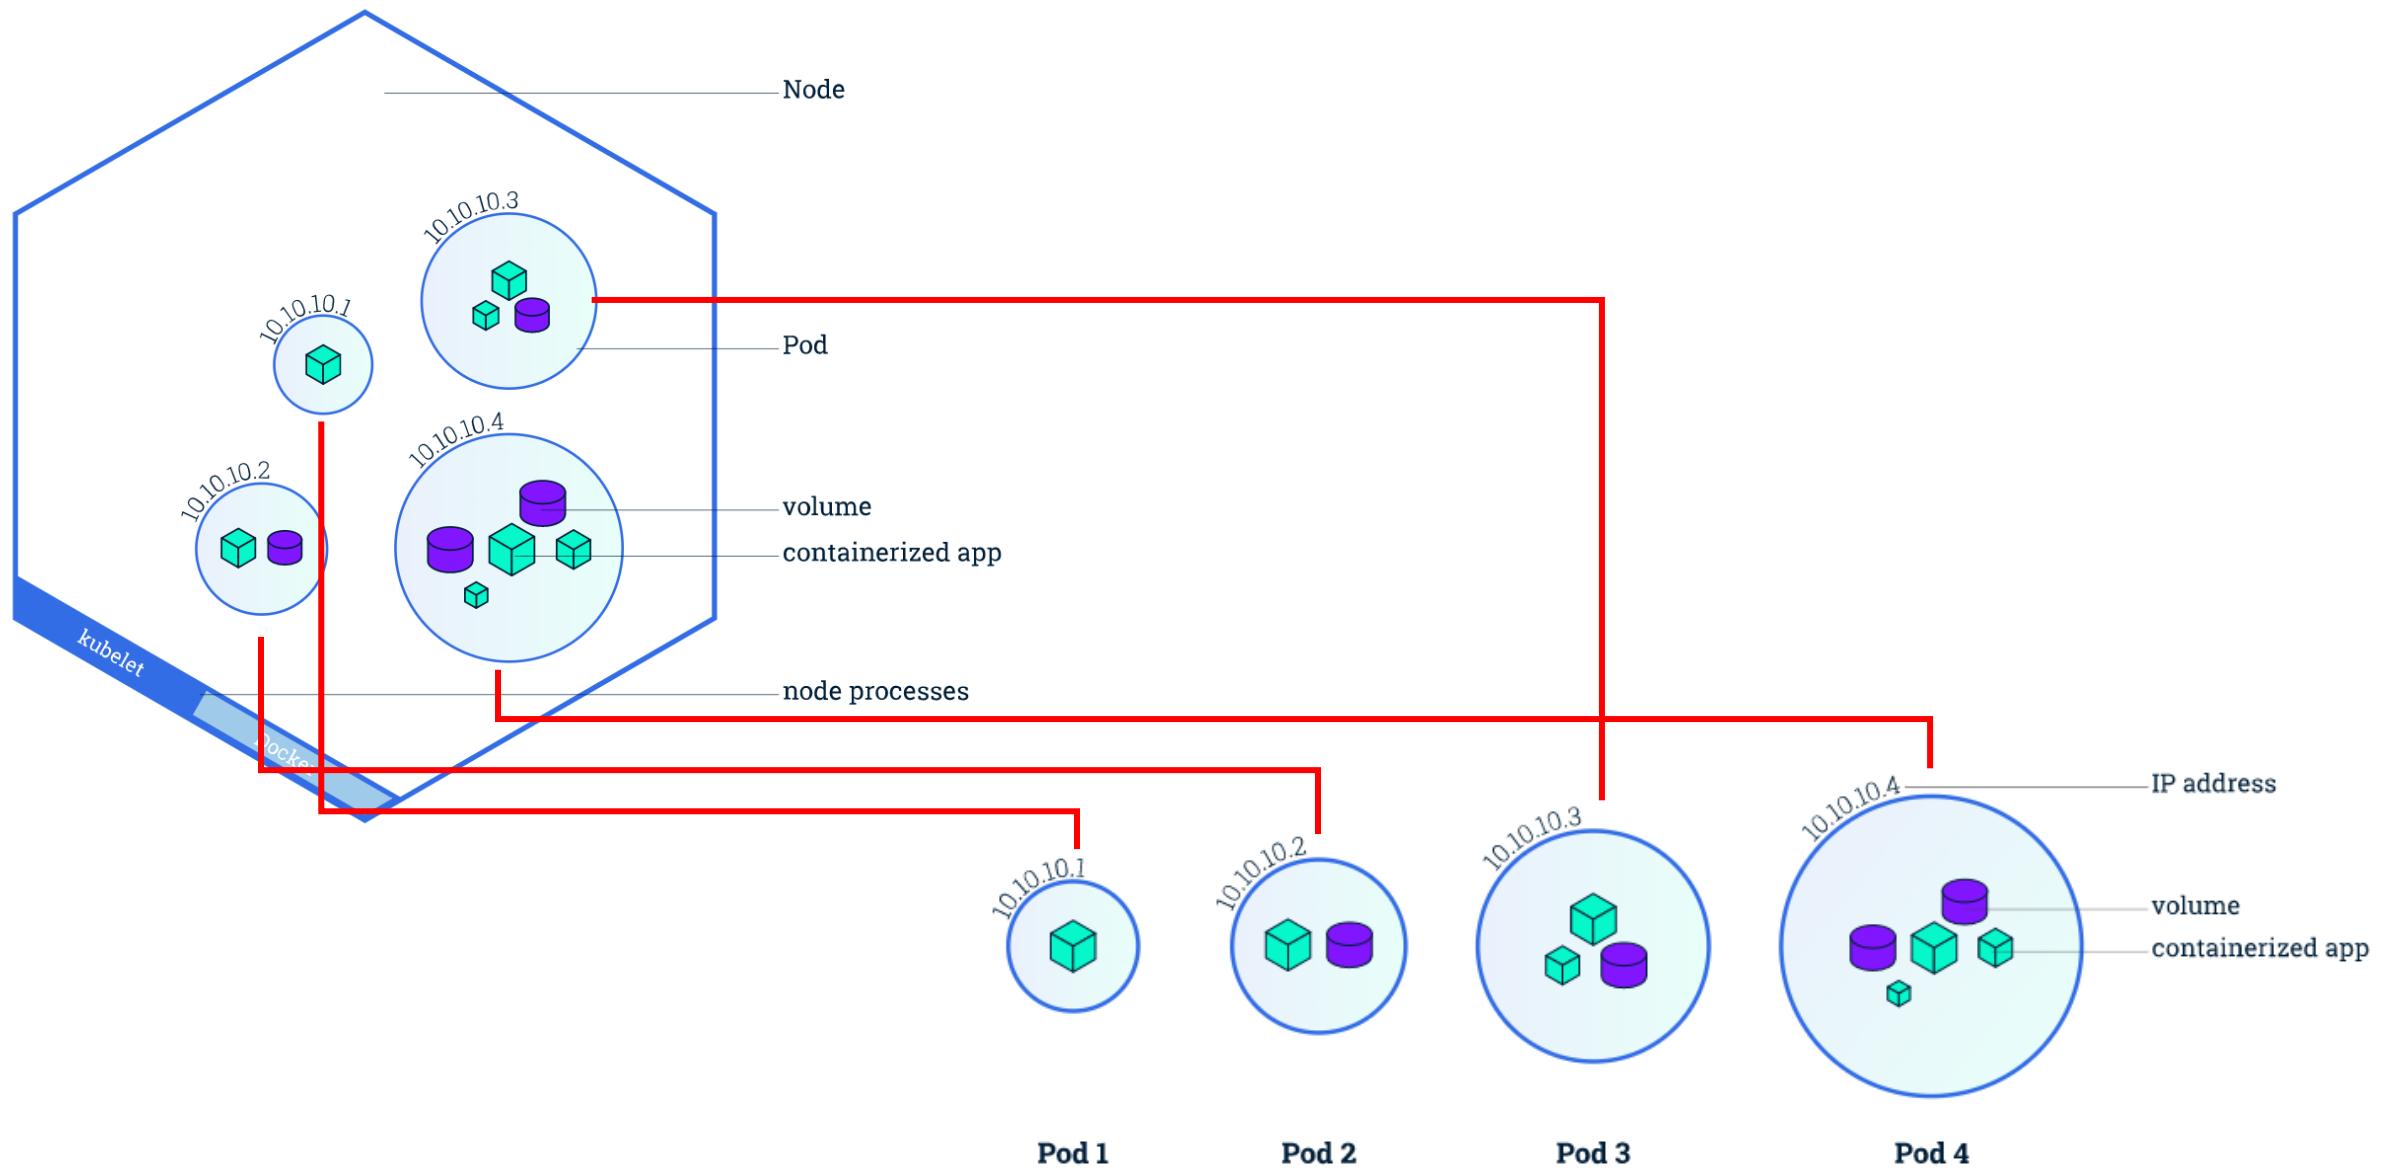
Kubernetes Architecture (Node)



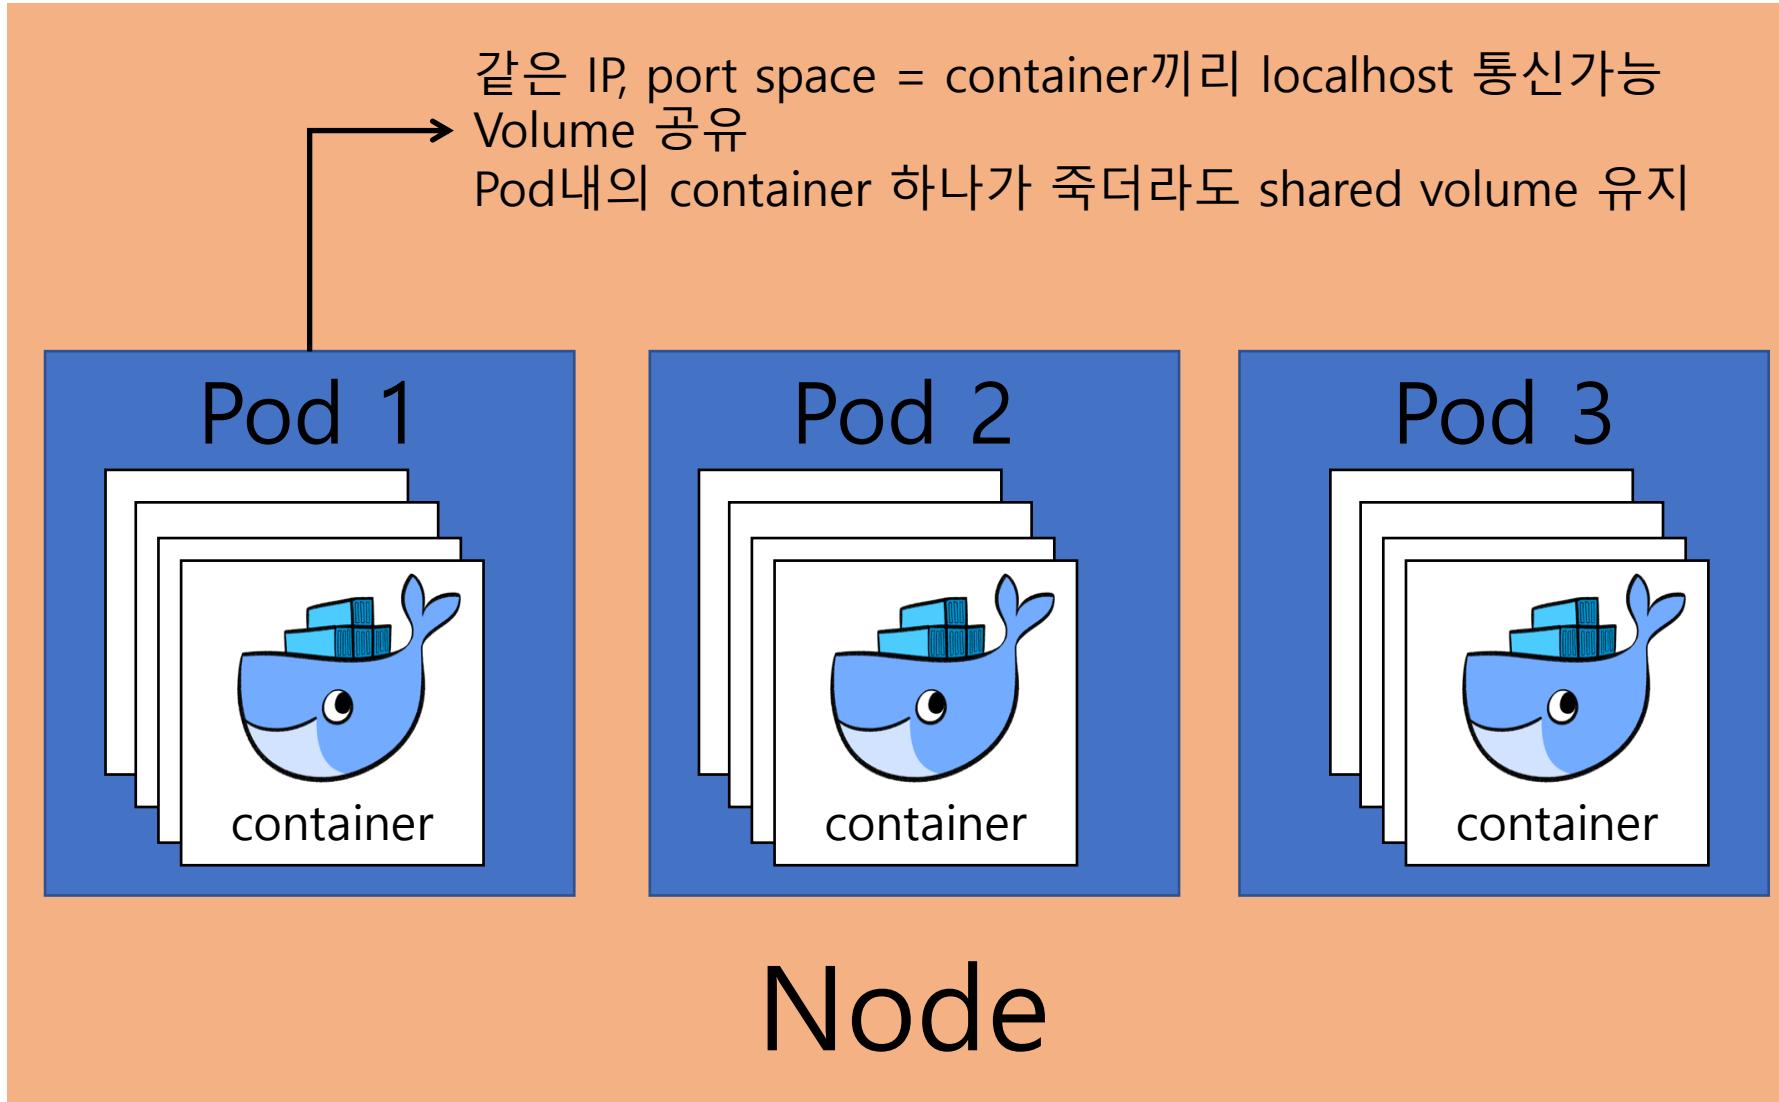
Kubernetes Object: Pod



Kubernetes Object: Pod

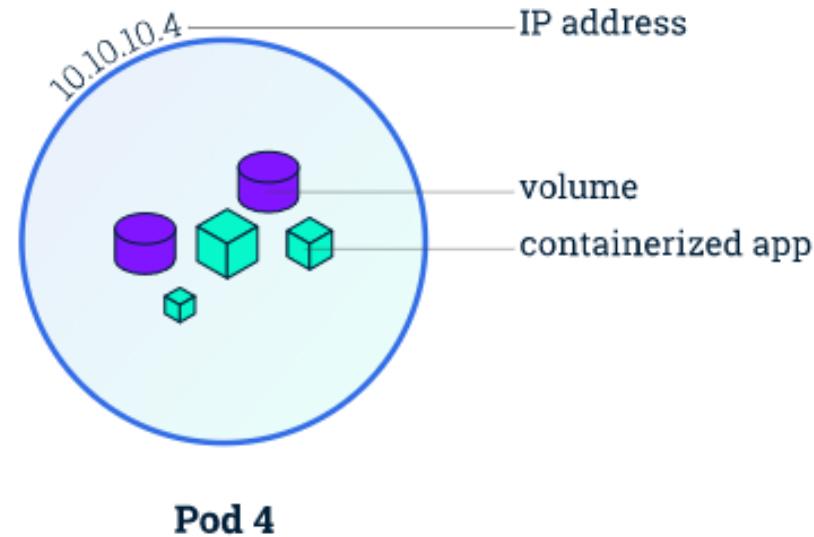


Kubernetes Object: Pod



Pod는 kubernetes의 작업 단위

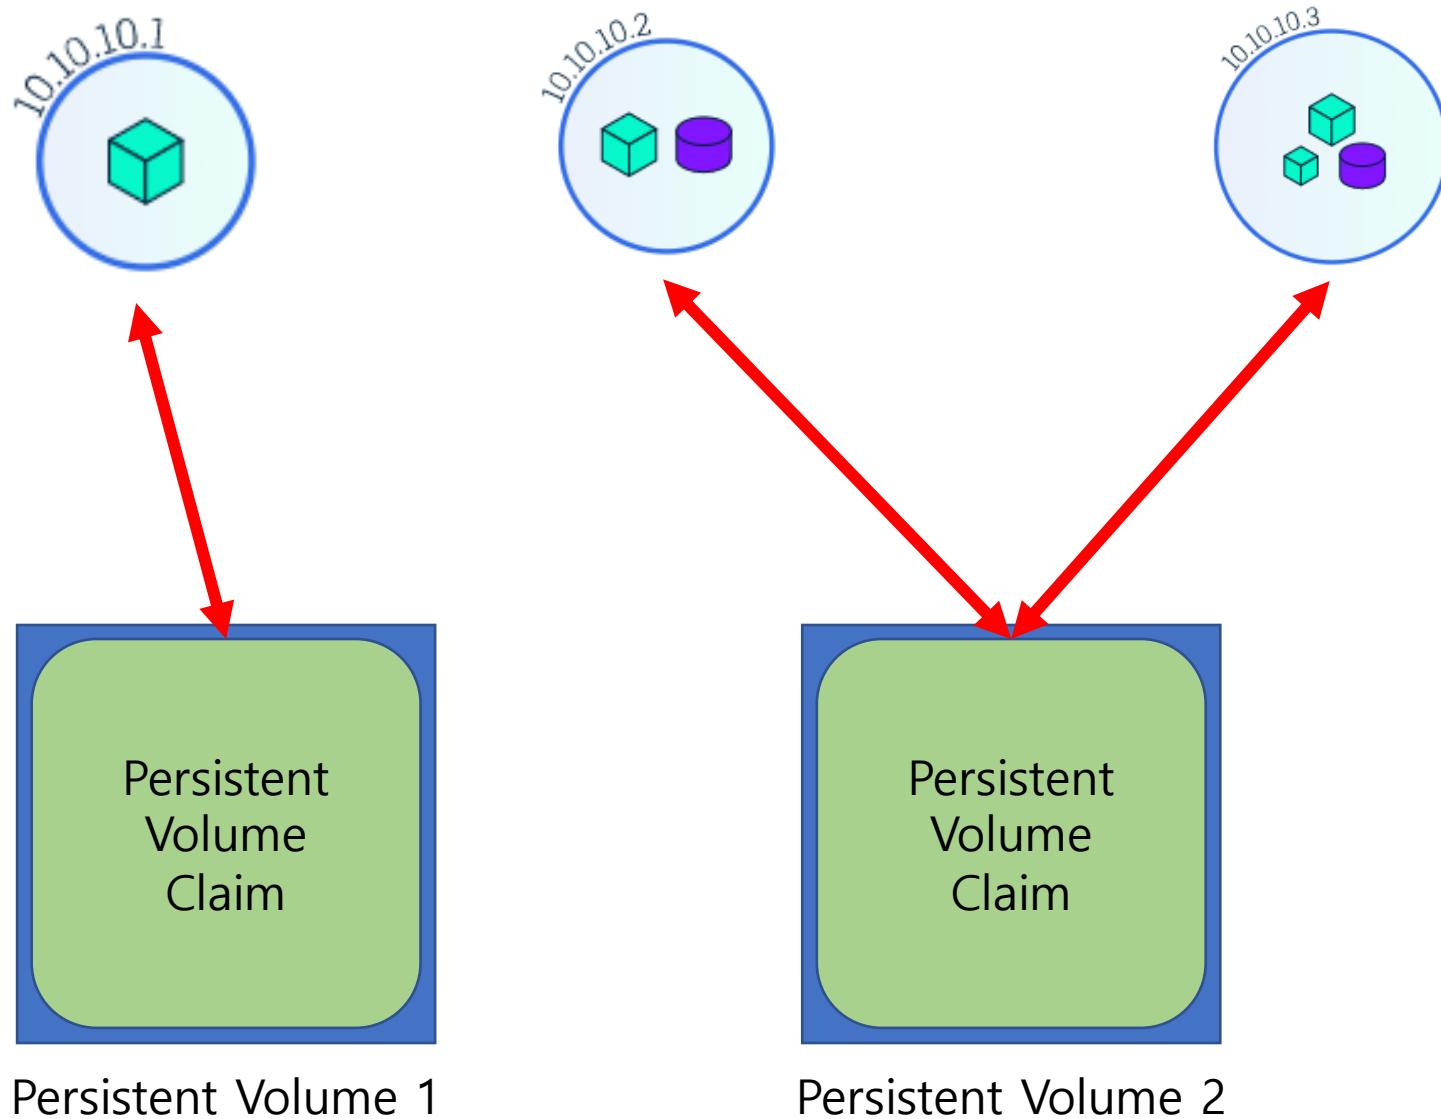
Kubernetes Object: Volume



stateless

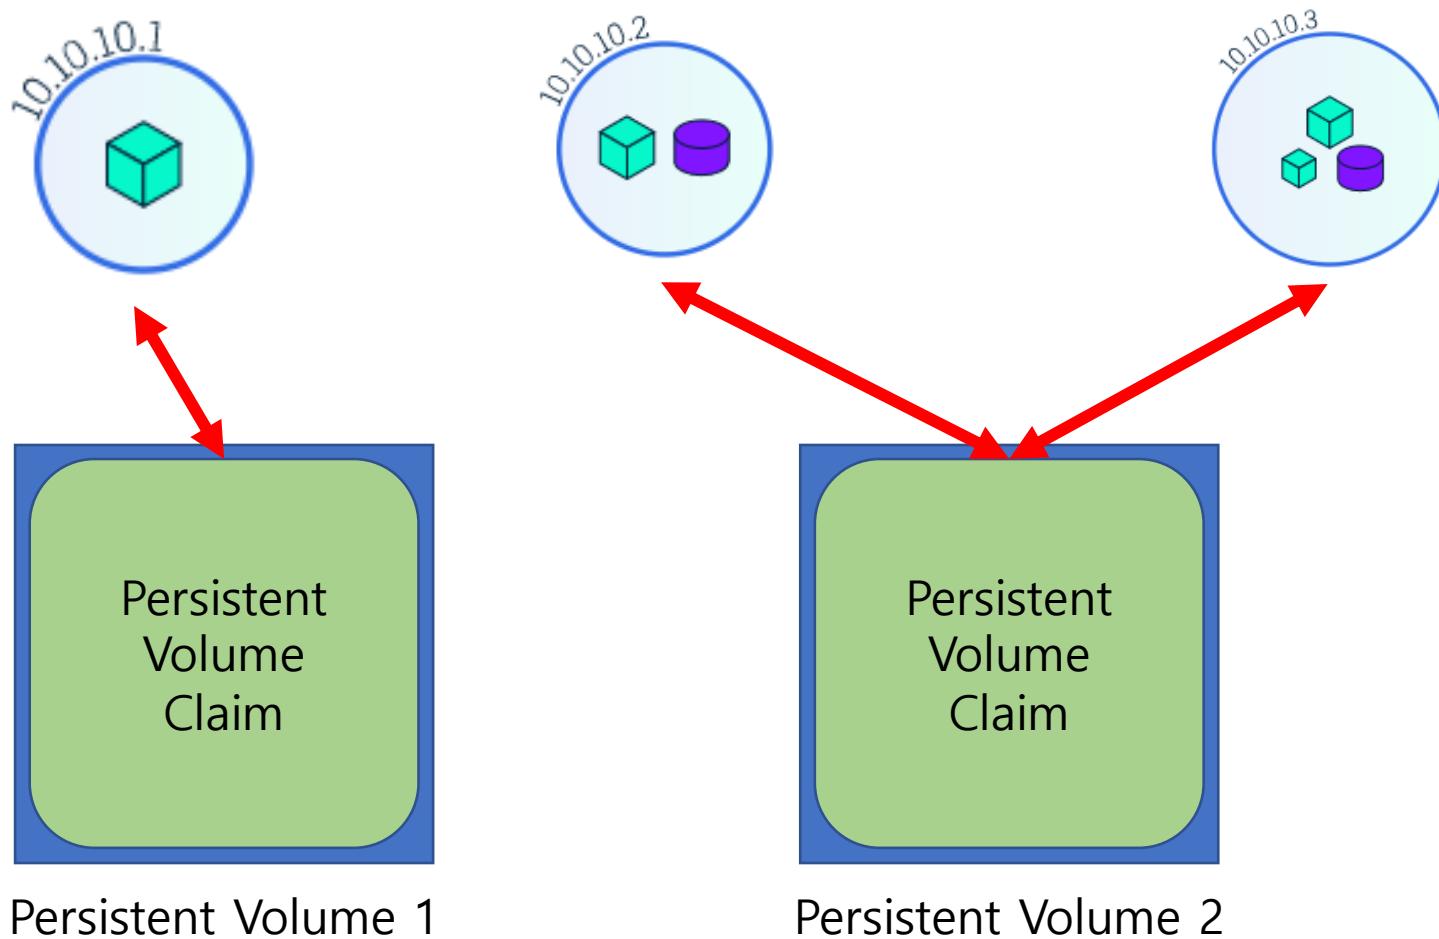
Pod는 상태저장이 되지 않음
즉 Pod가 사라지면 Pod의 volume도 사라짐

Kubernetes Object: Volume



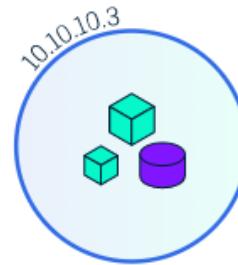
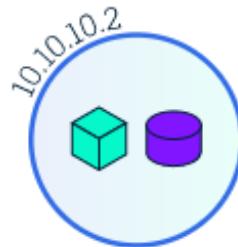
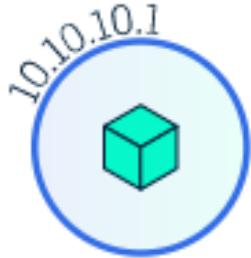
Kubernetes Object: Volume

정적 프로비저닝



Kubernetes Object: Volume

정적 프로비저닝



미리 PV를 만들어 놓음



Persistent Volume 1



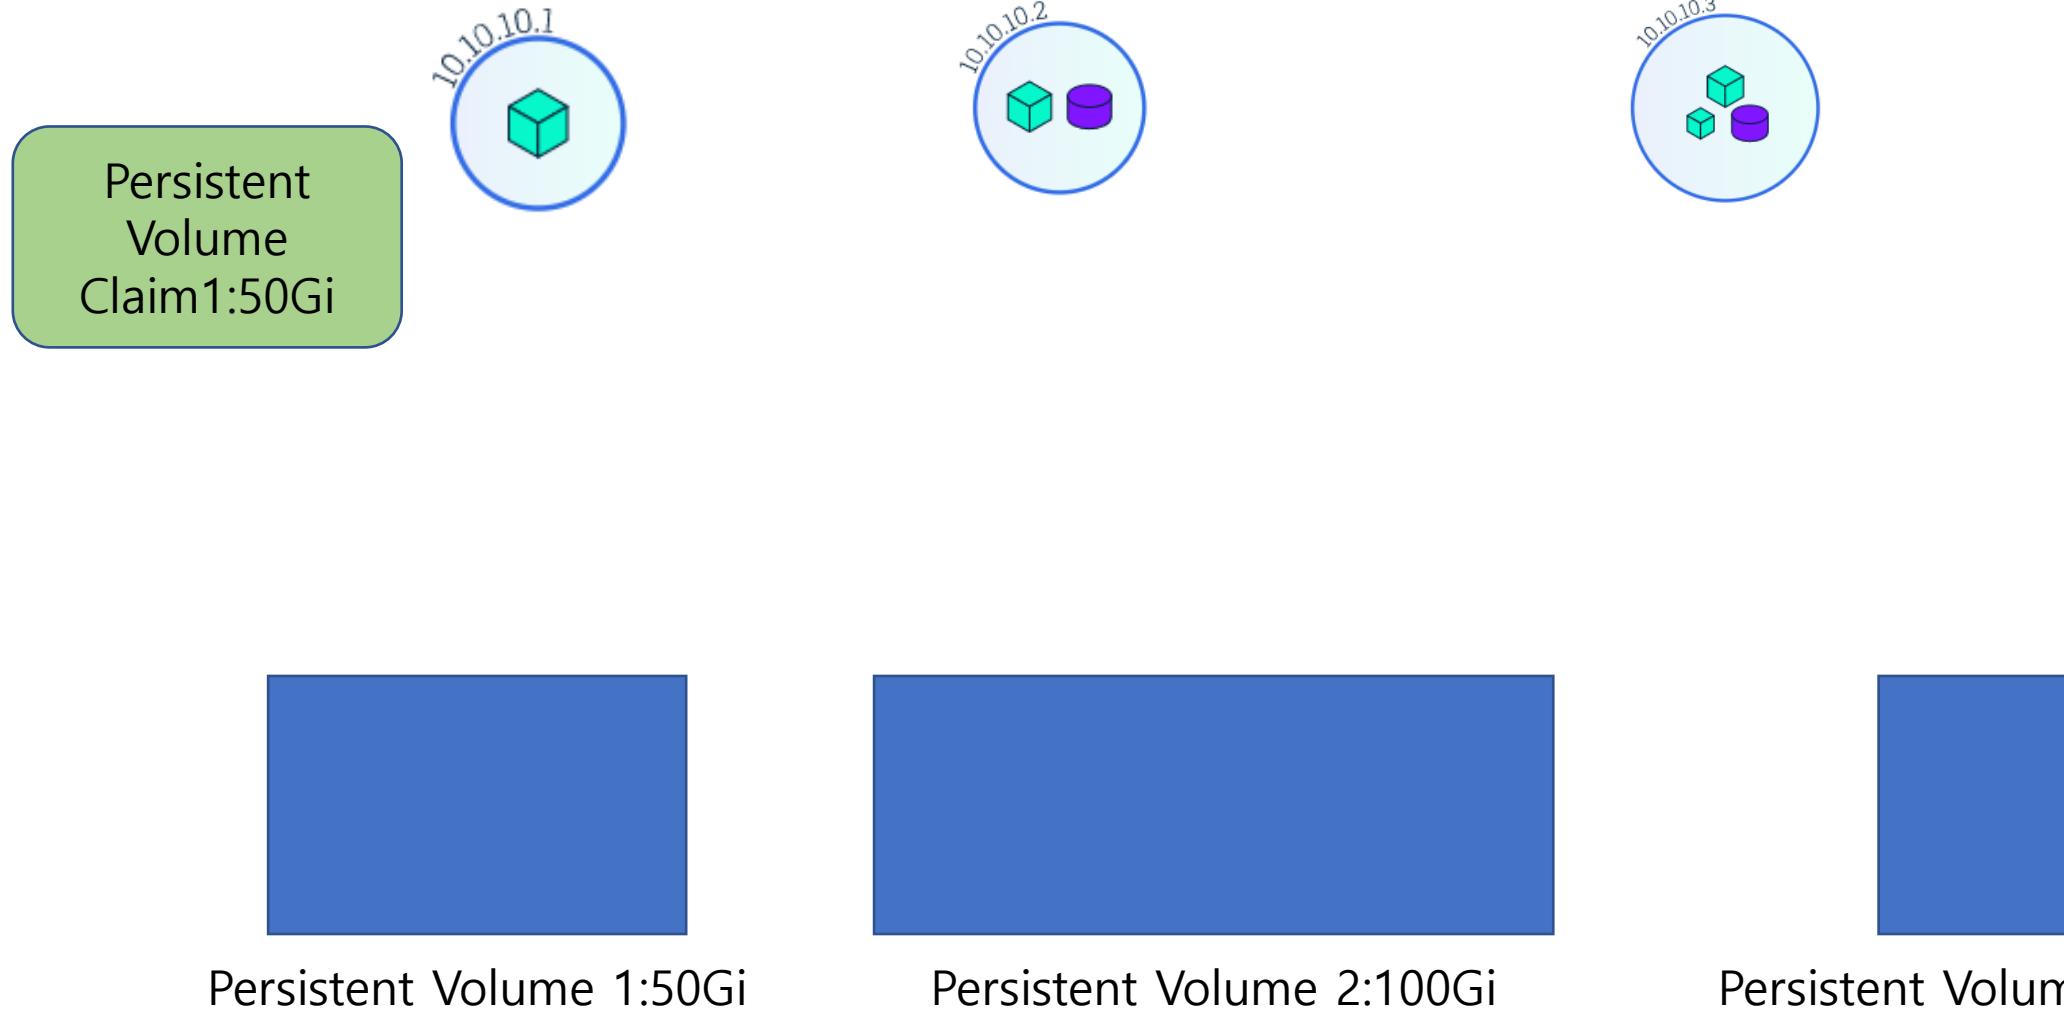
Persistent Volume 2



Persistent Volume 3

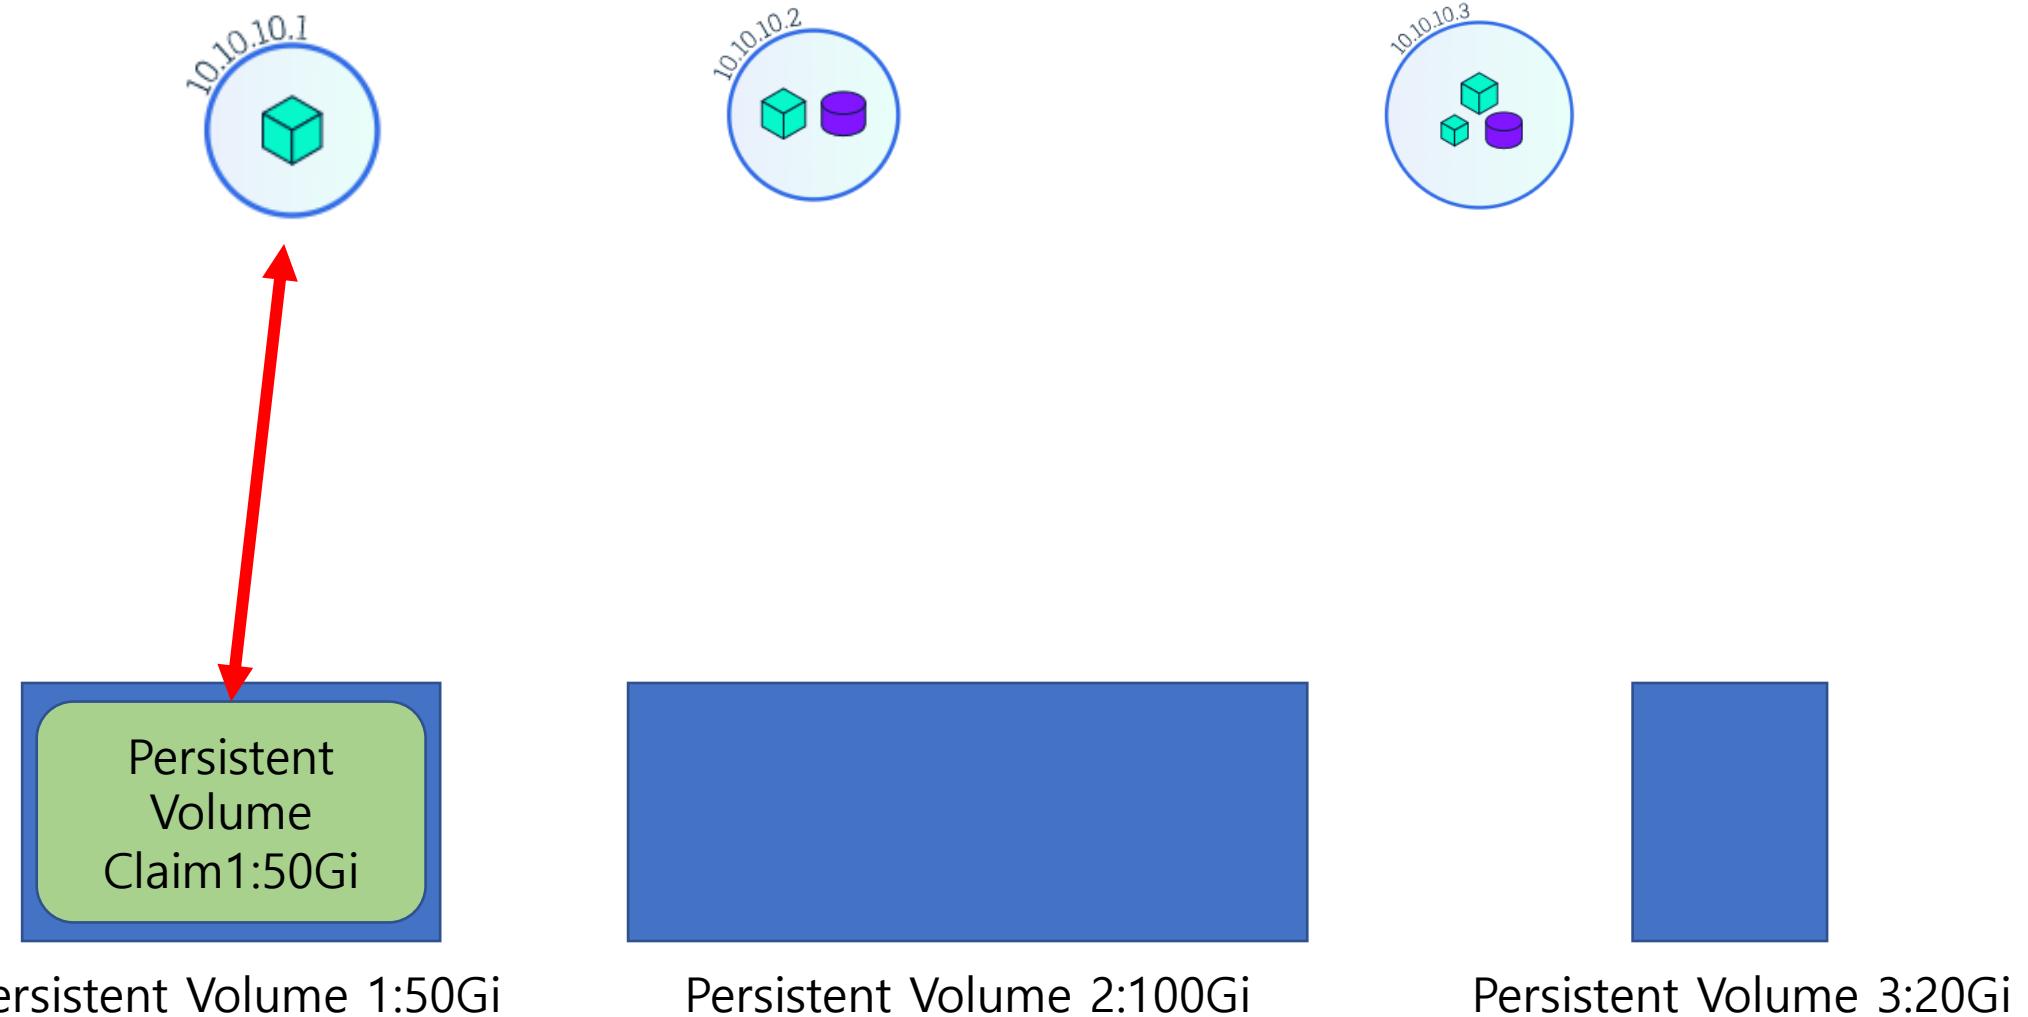
Kubernetes Object: Volume

정적 프로비저닝



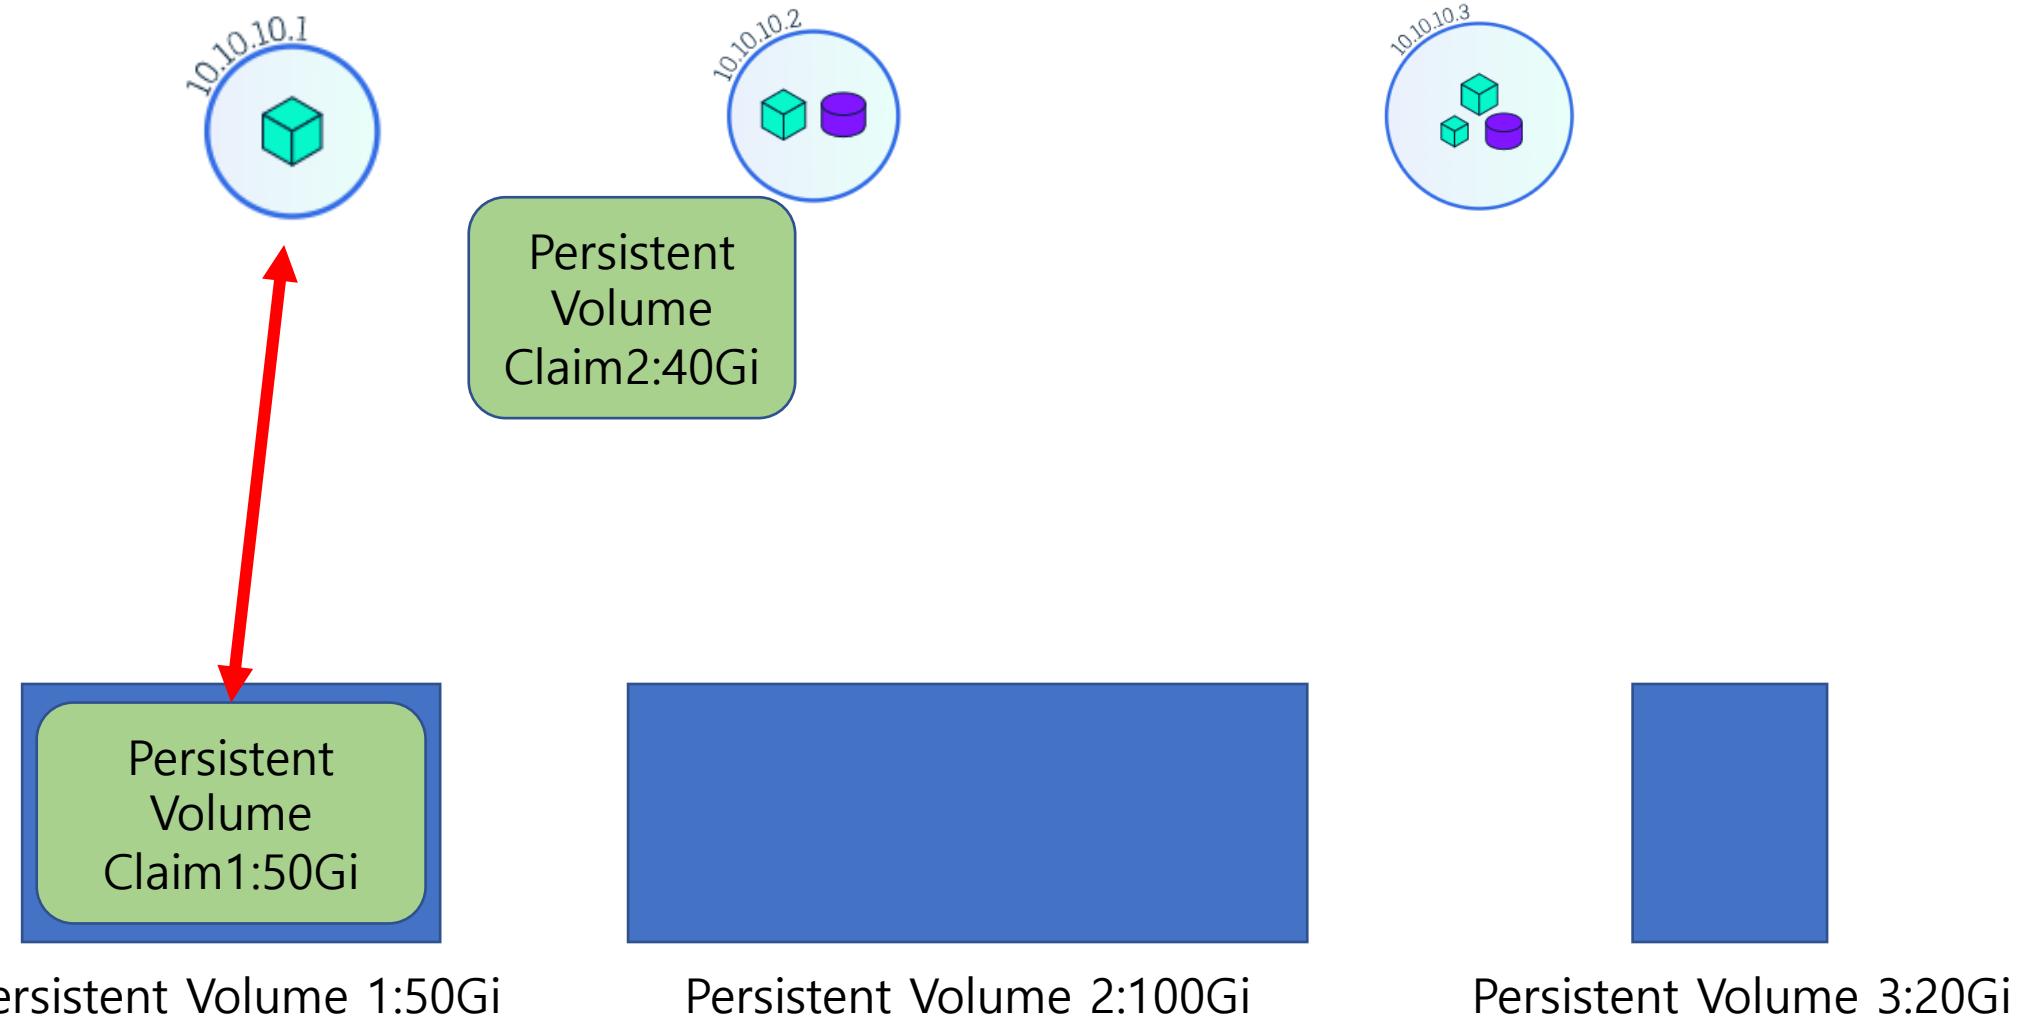
Kubernetes Object: Volume

정적 프로비저닝



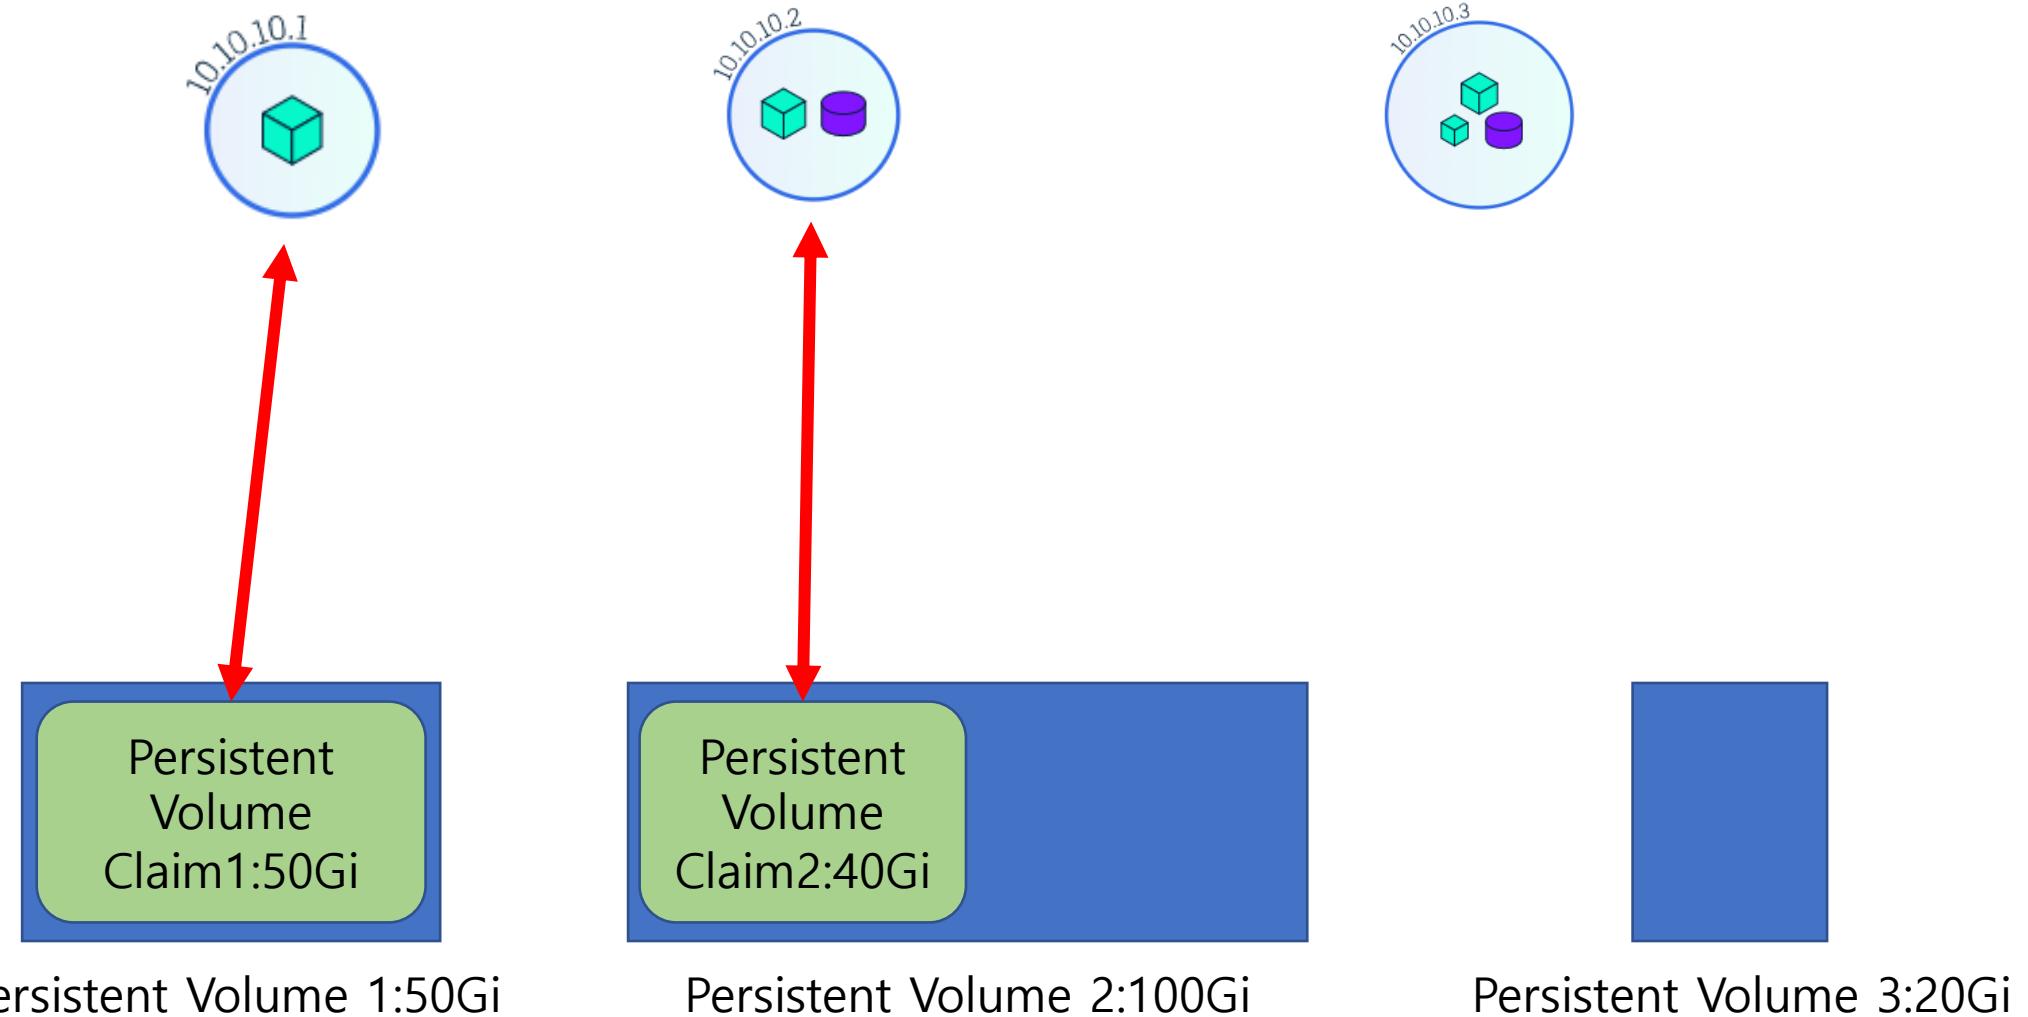
Kubernetes Object: Volume

정적 프로비저닝



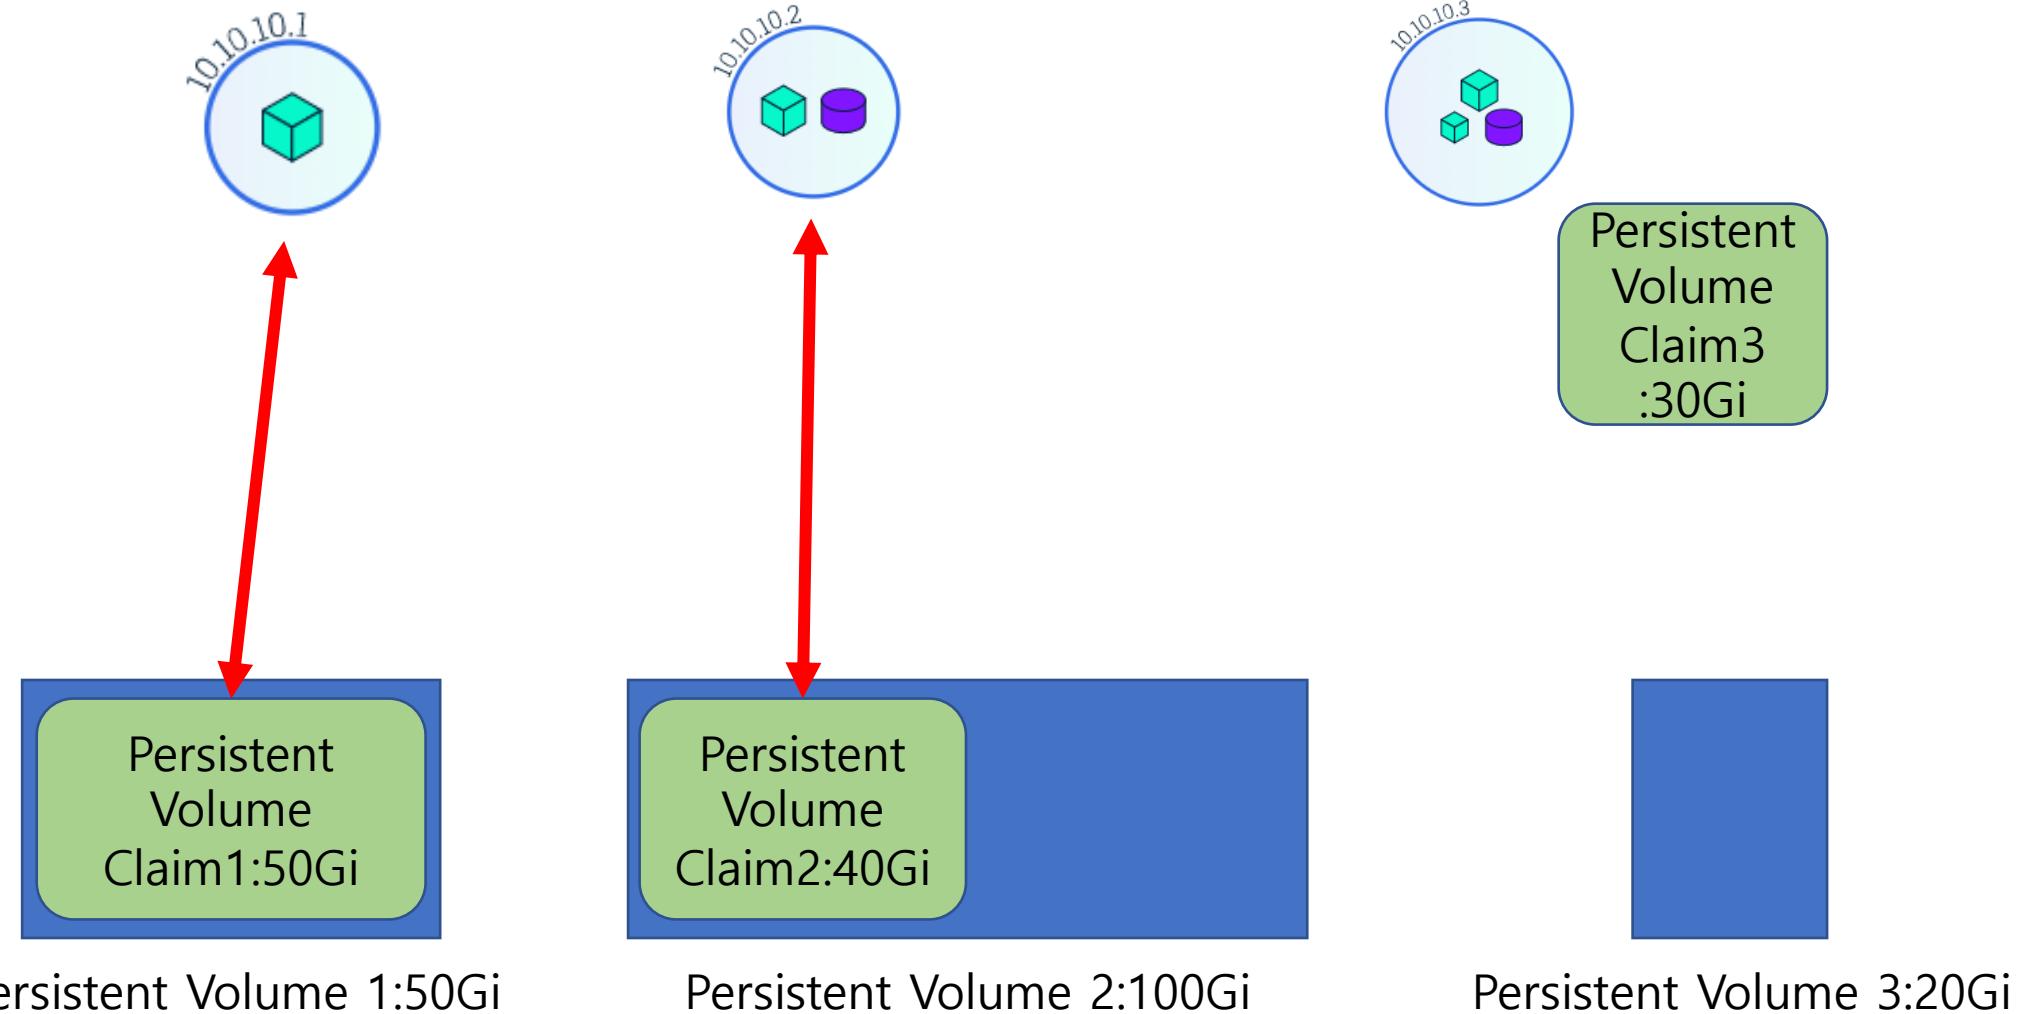
Kubernetes Object: Volume

정적 프로비저닝



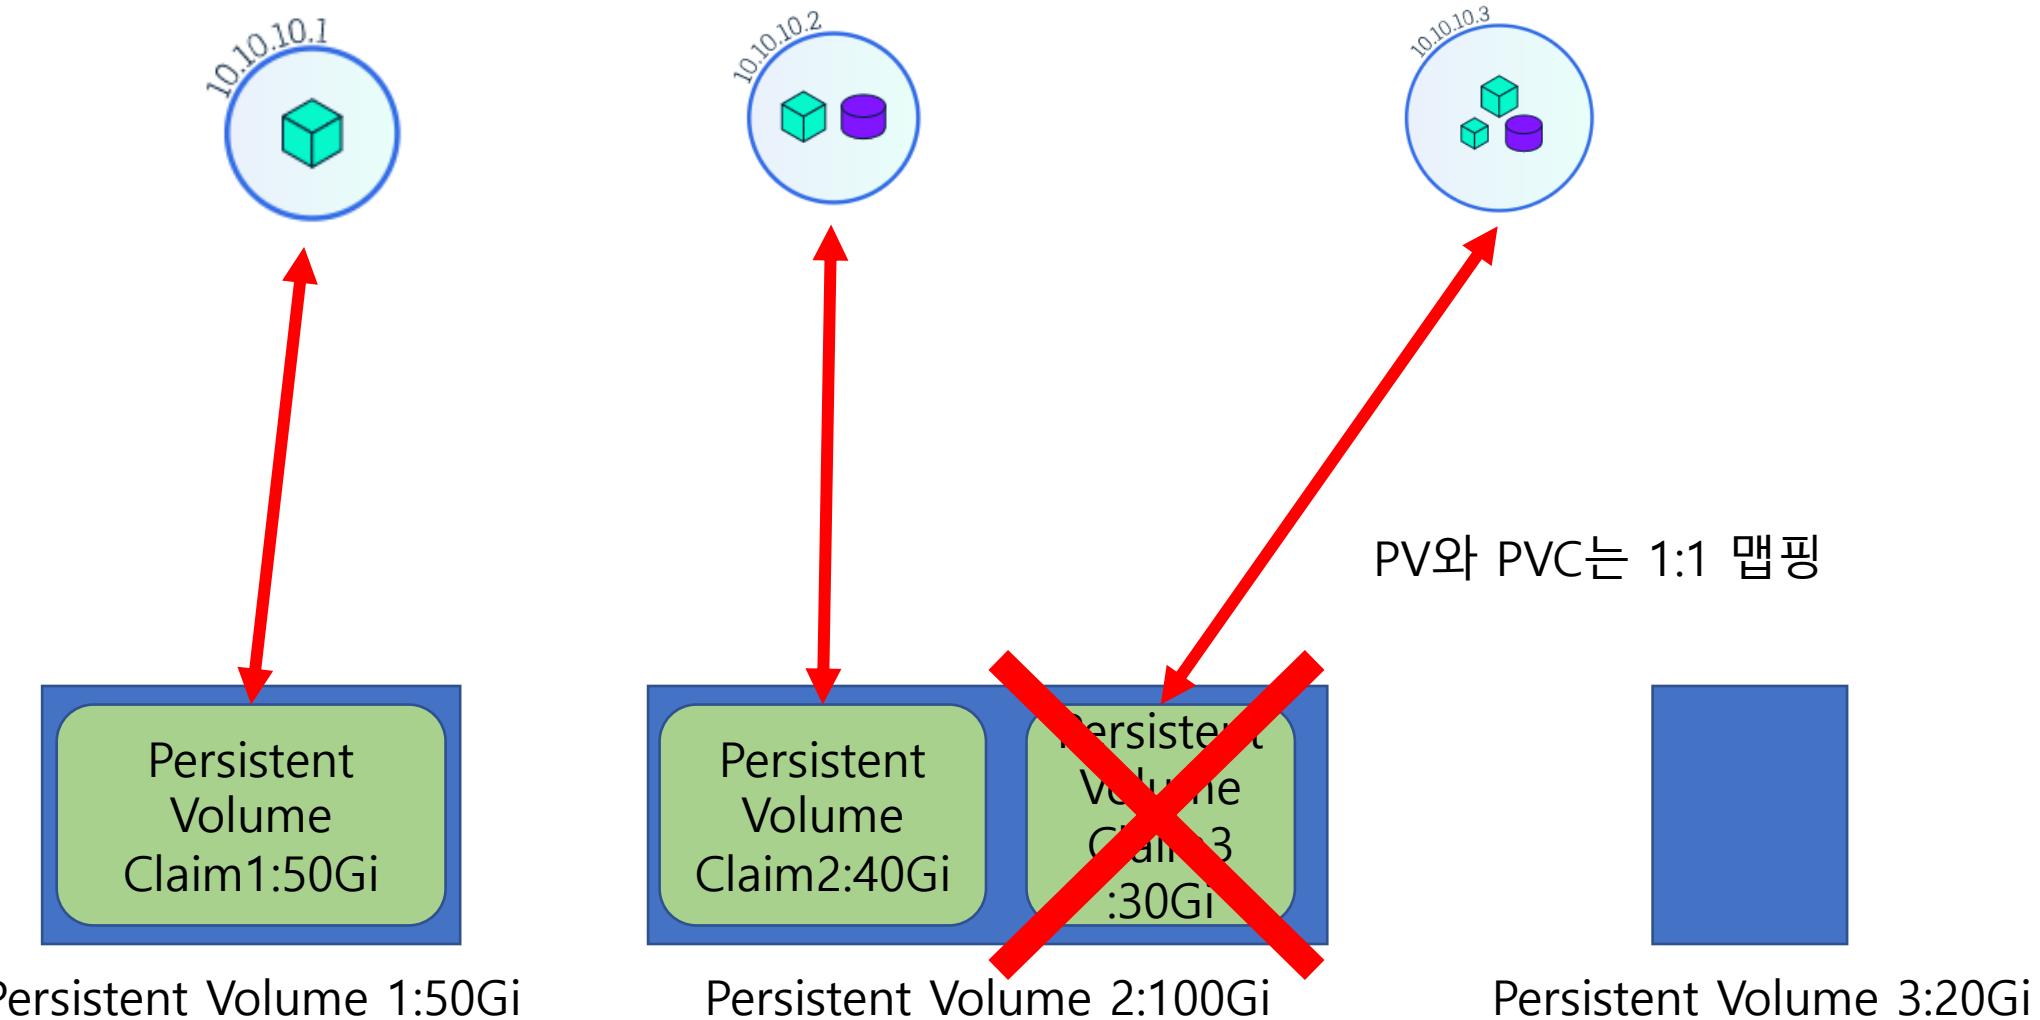
Kubernetes Object: Volume

정적 프로비저닝



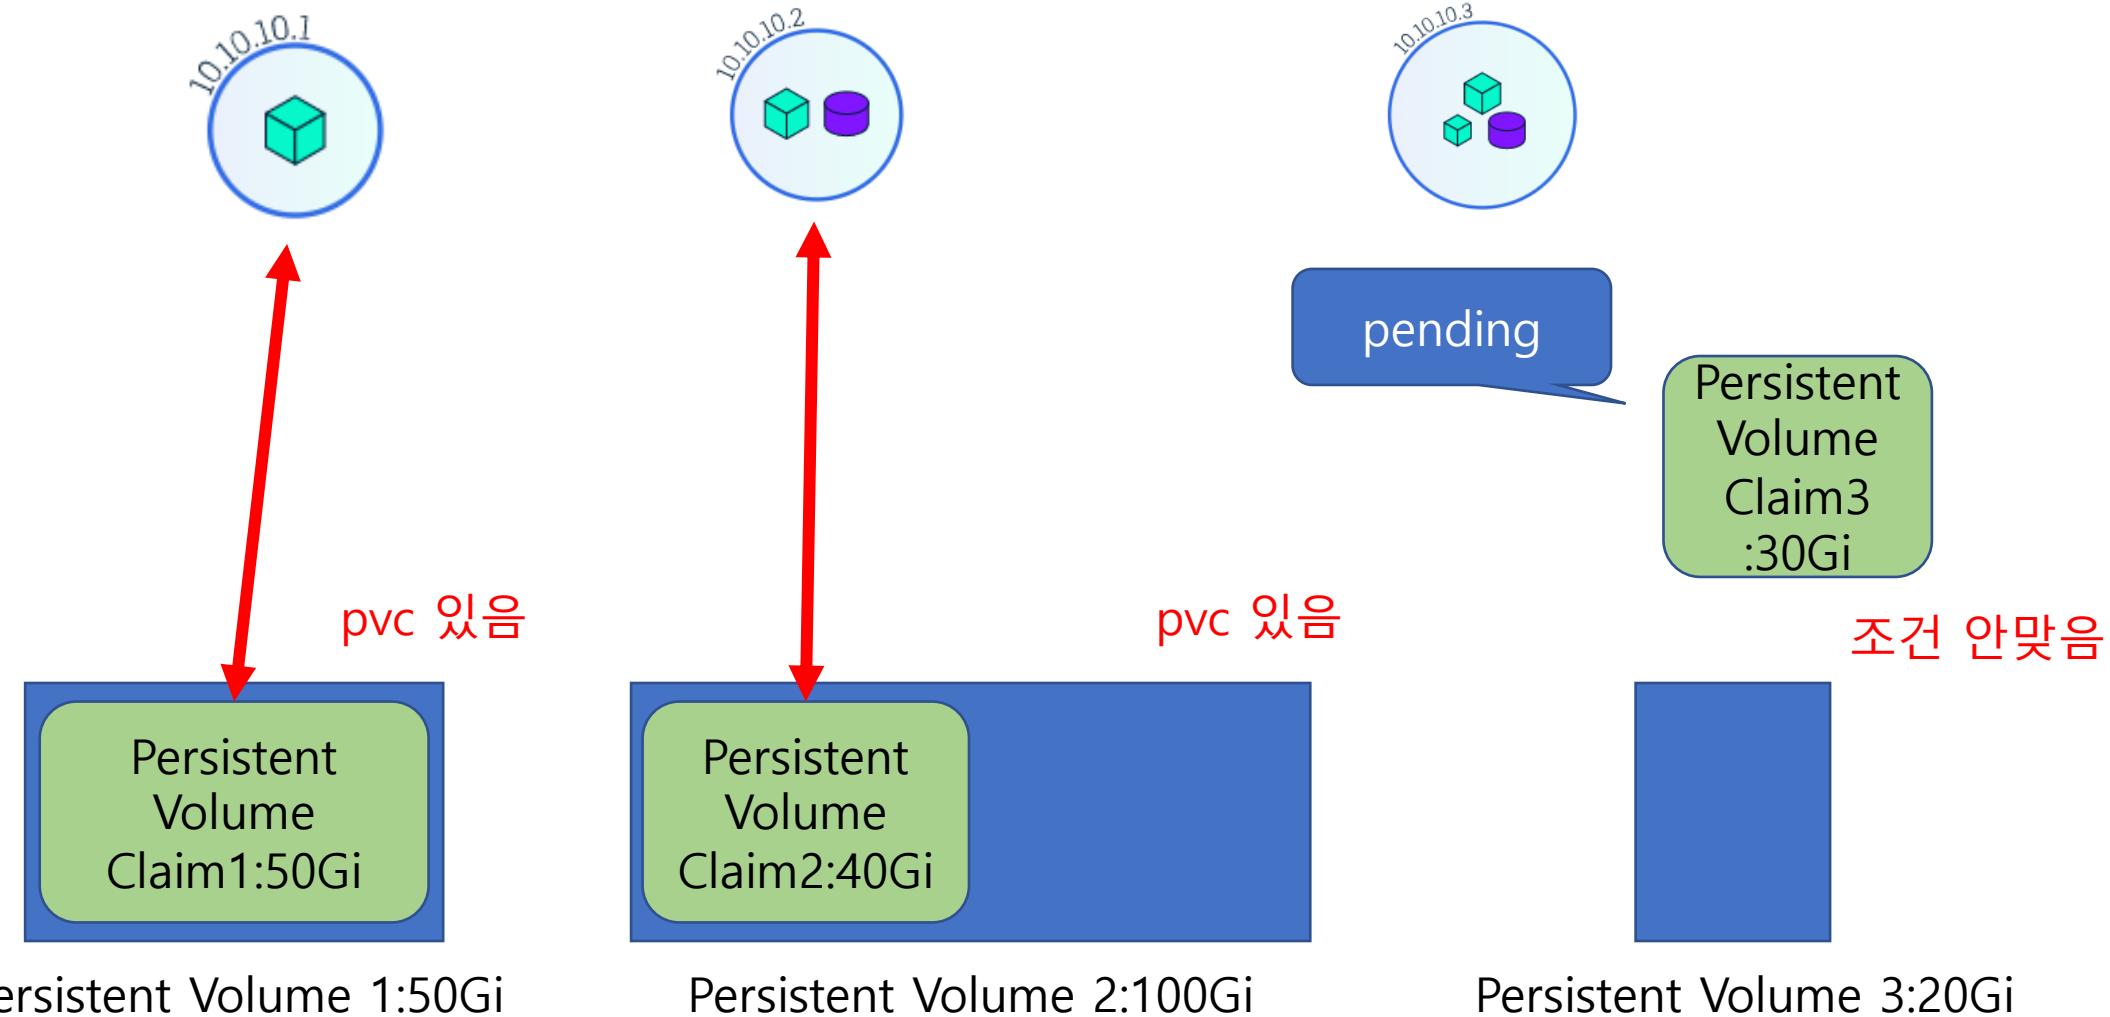
Kubernetes Object: Volume

정적 프로비저닝



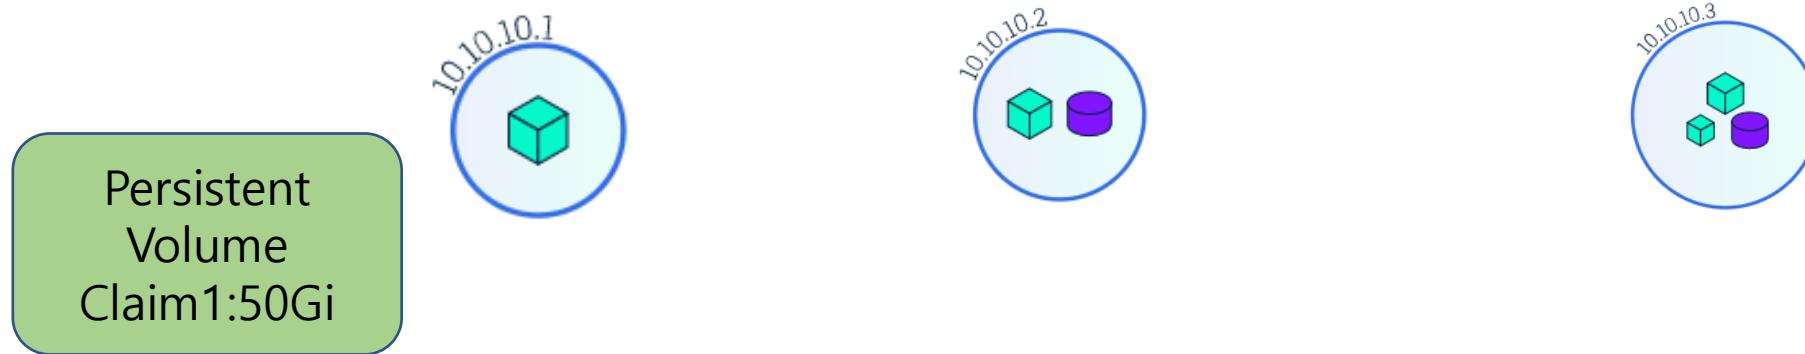
Kubernetes Object: Volume

정적 프로비저닝



Kubernetes Object: Volume

동적 프로비저닝



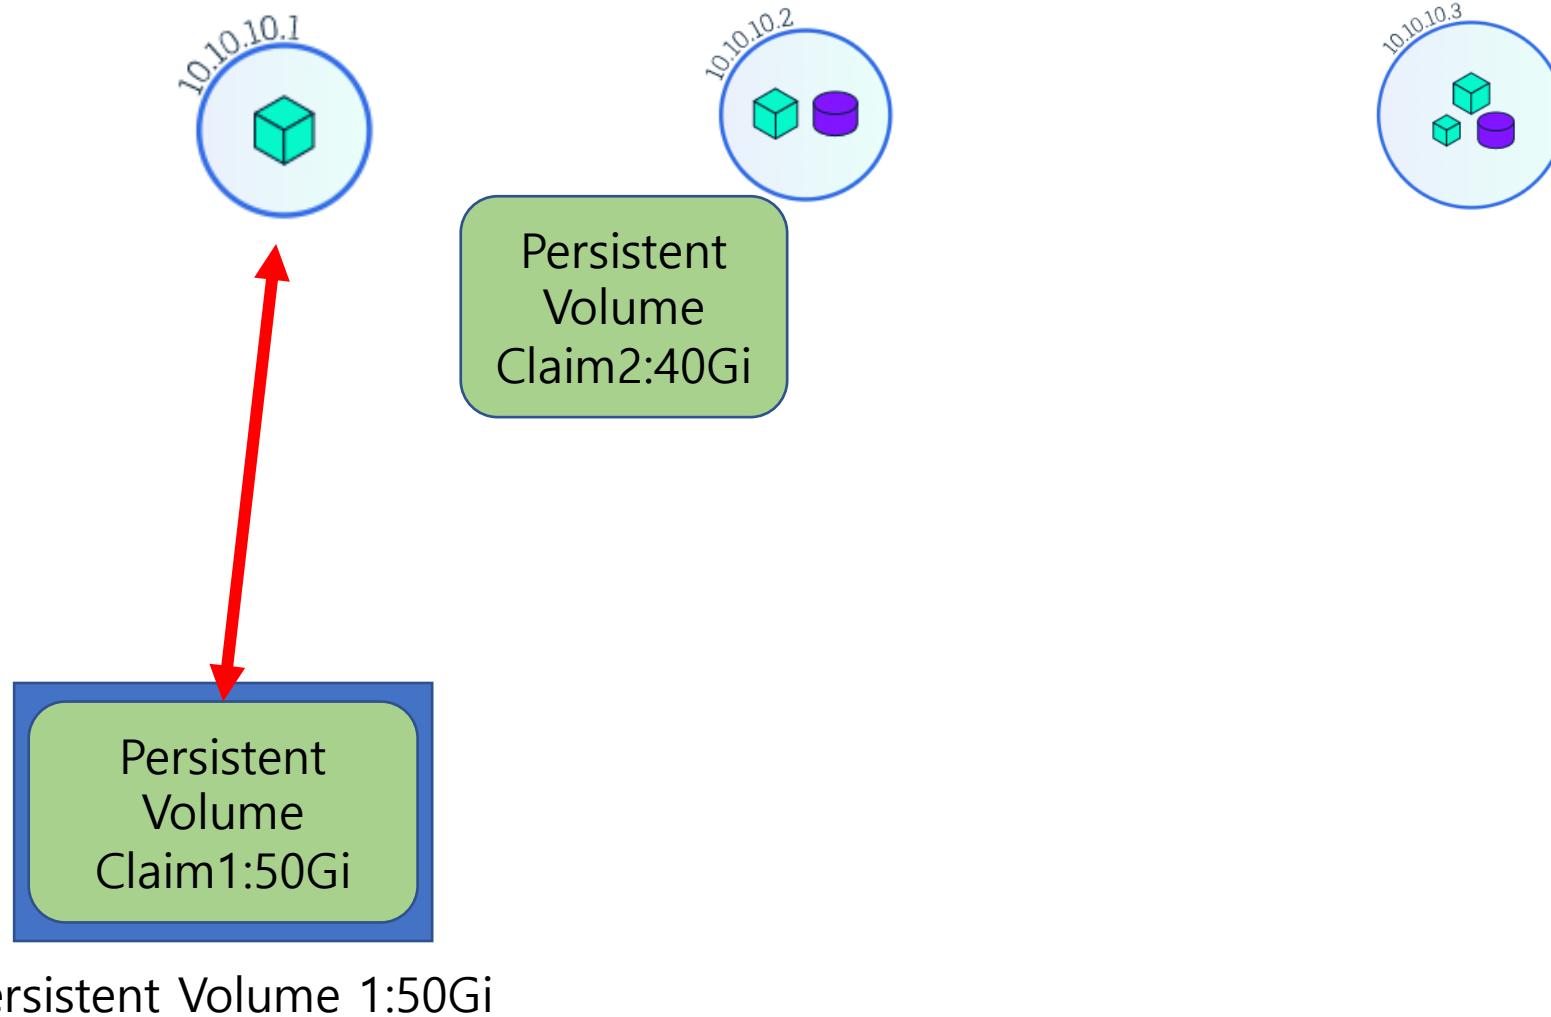
Kubernetes Object: Volume

동적 프로비저닝



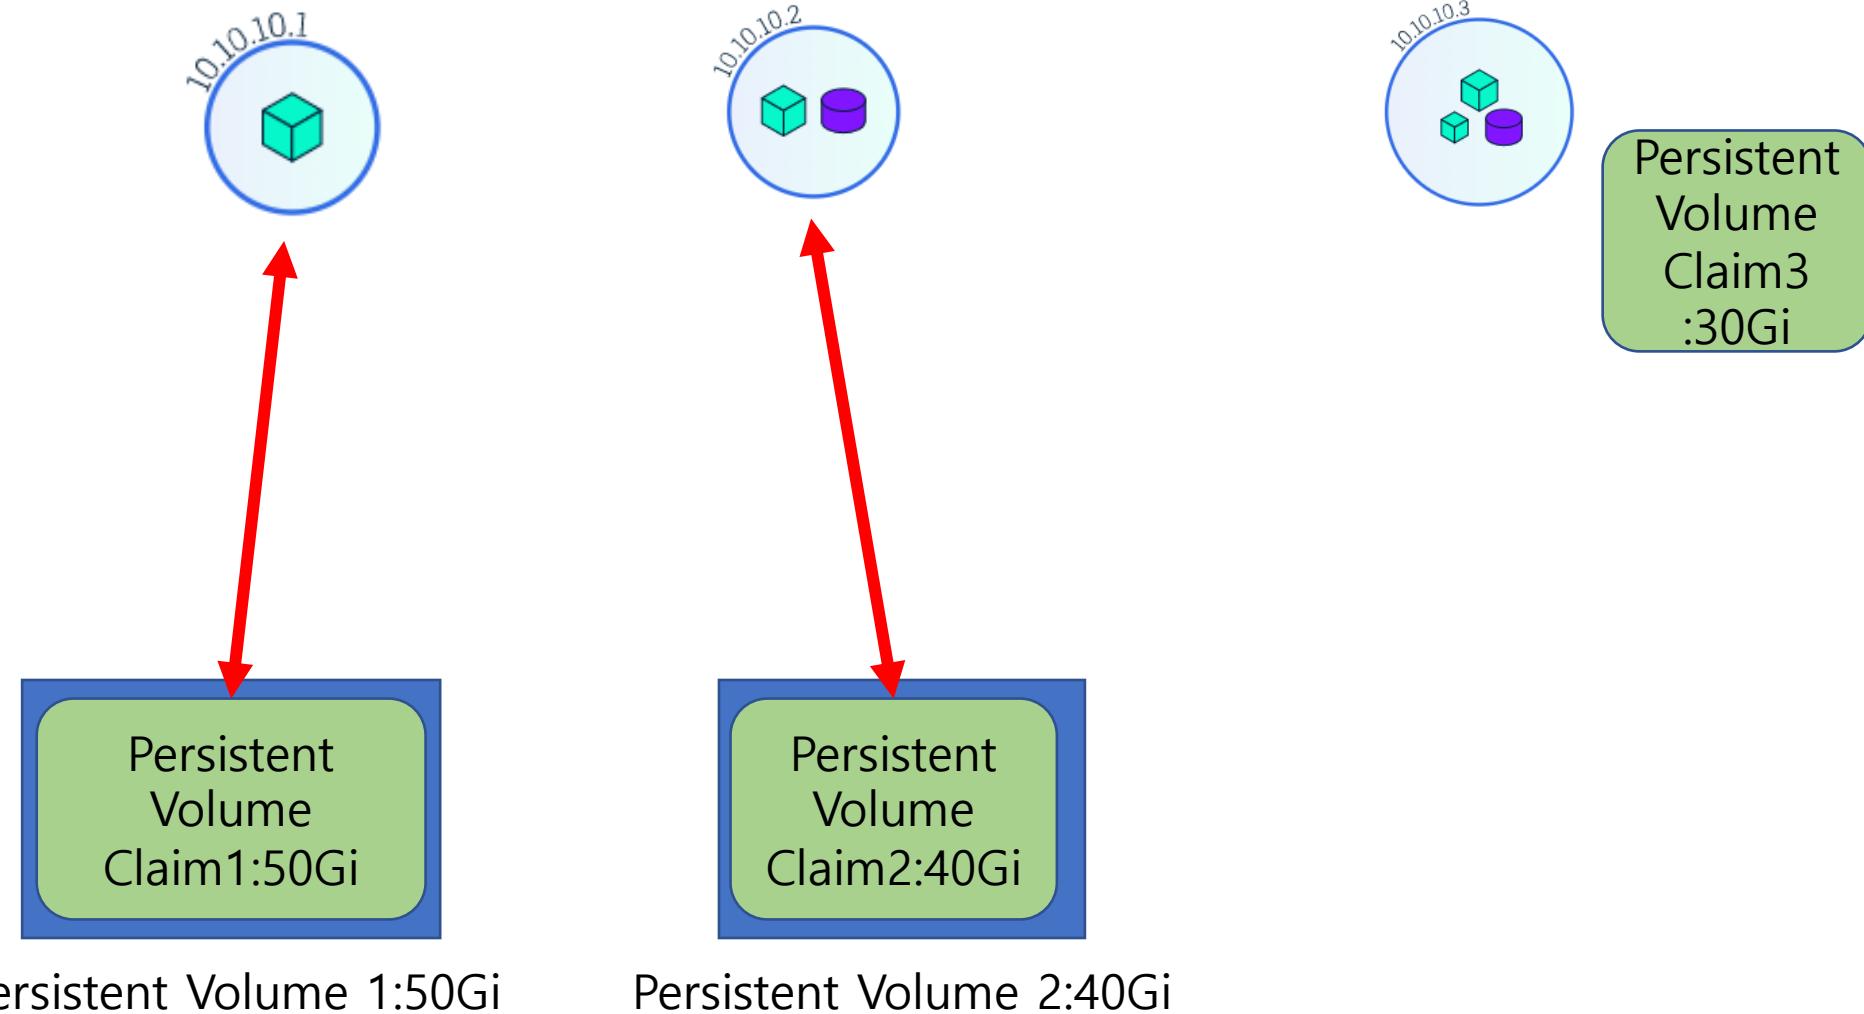
Kubernetes Object: Volume

동적 프로비저닝



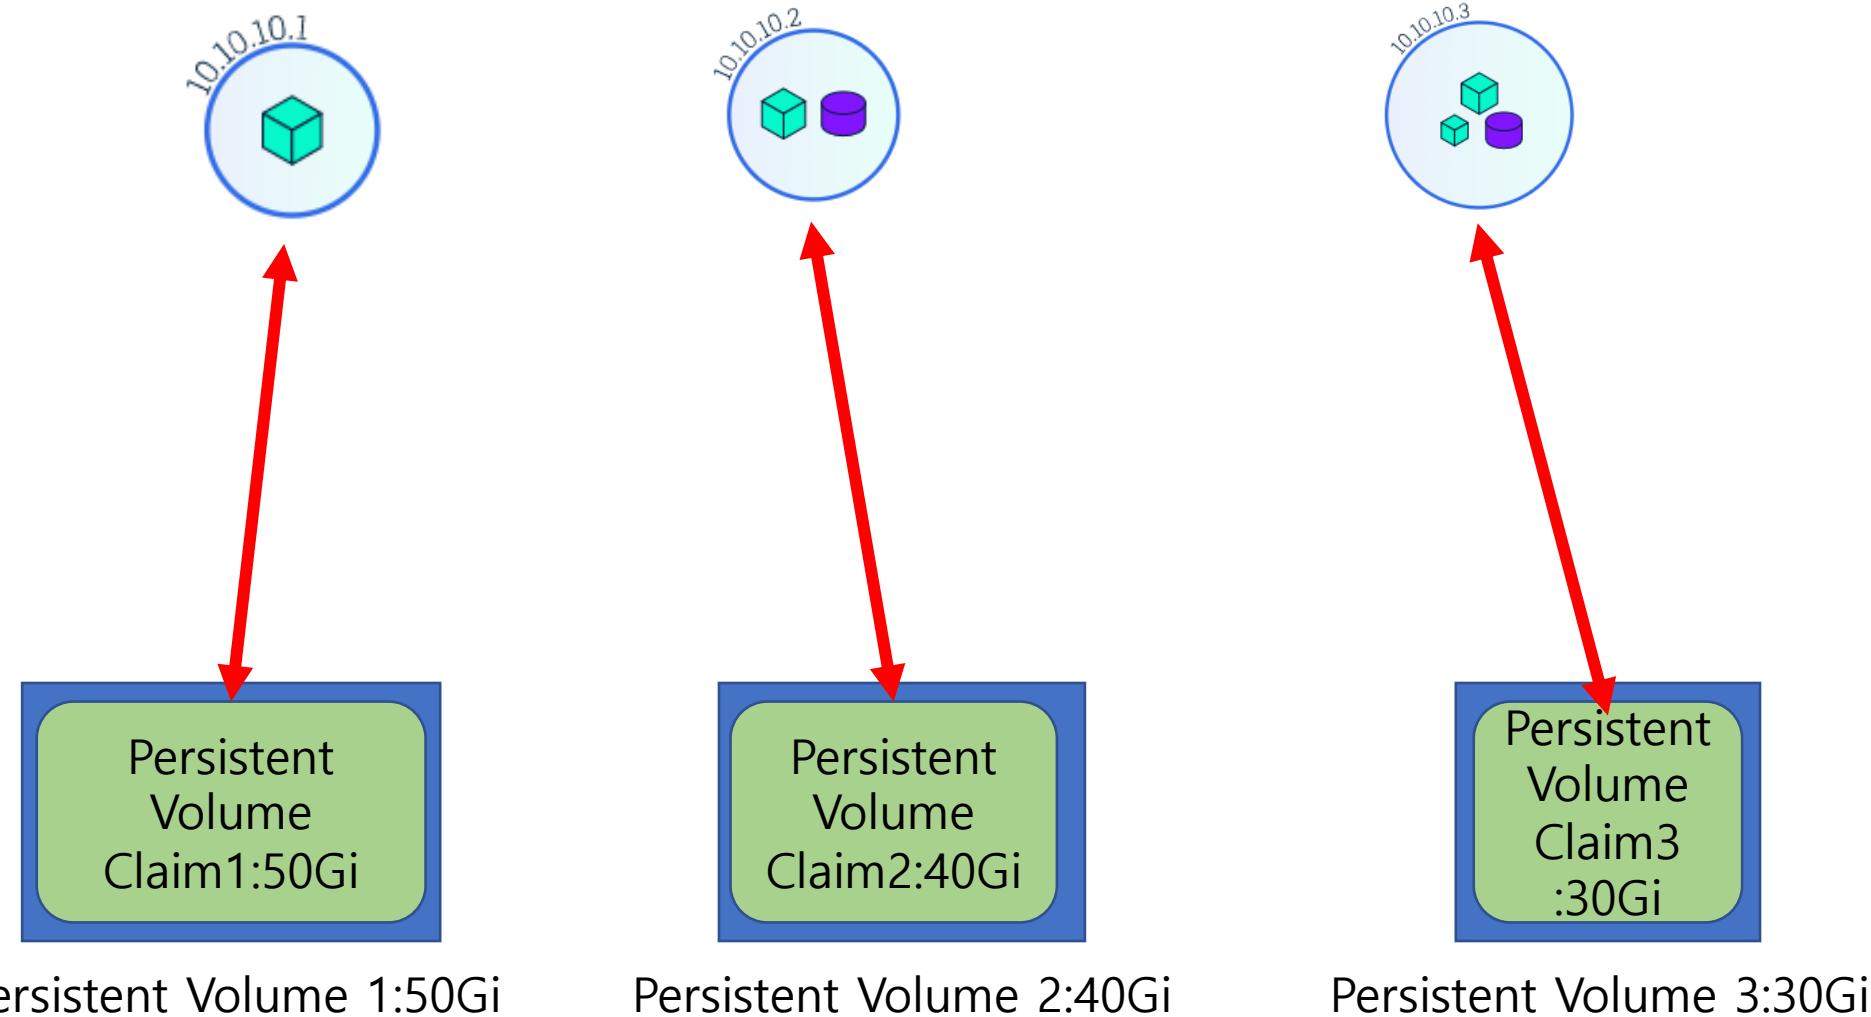
Kubernetes Object: Volume

동적 프로비저닝



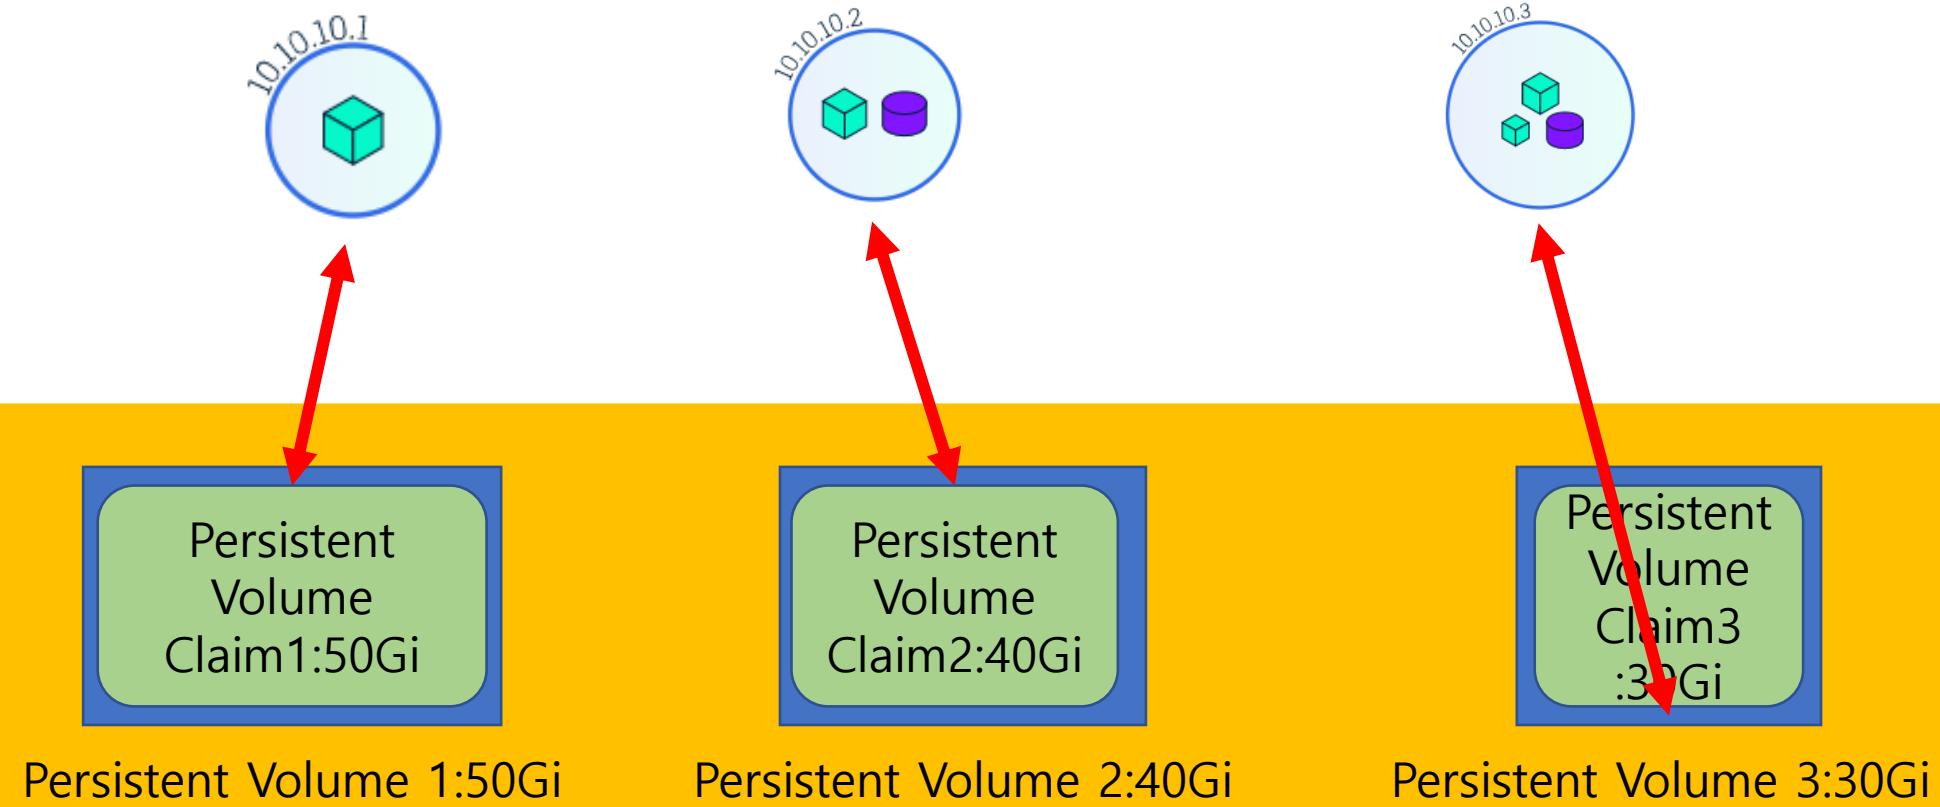
Kubernetes Object: Volume

동적 프로비저닝



Kubernetes Object: Volume

동적 프로비저닝



StorageClass

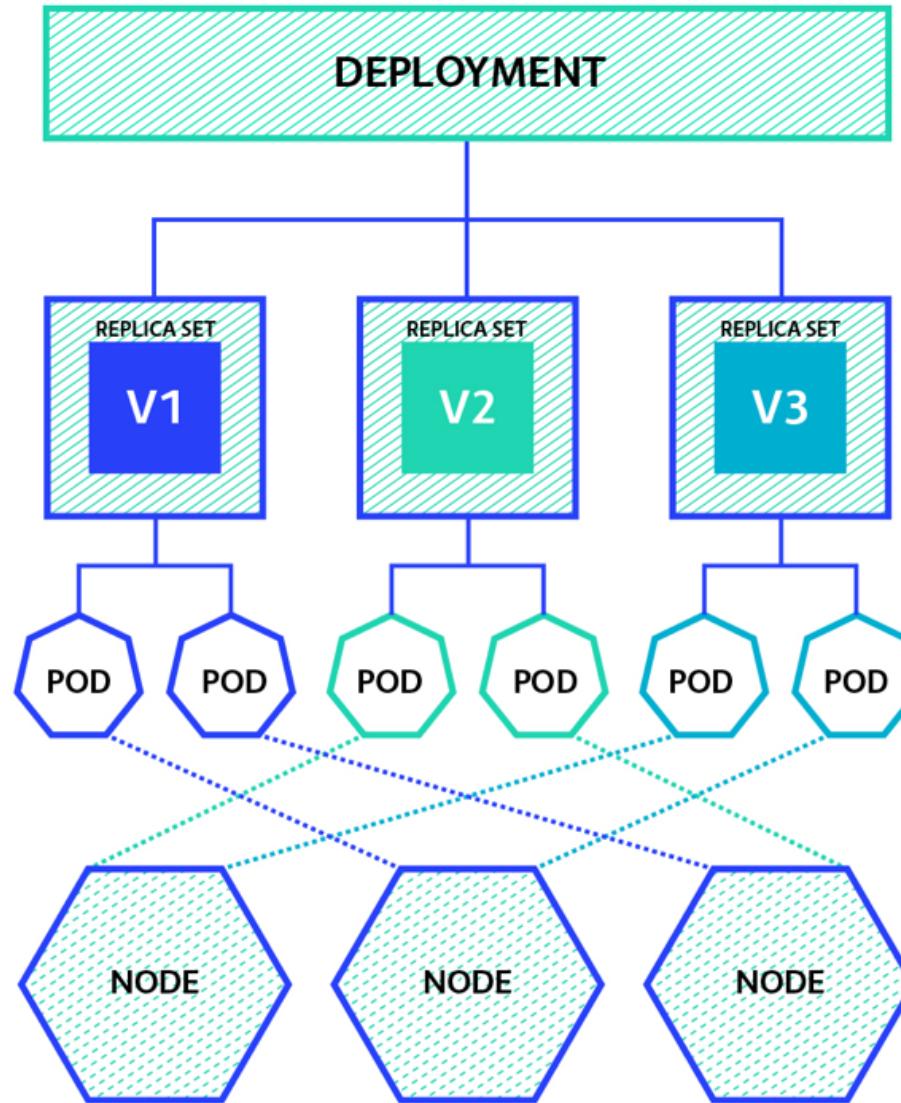
StorageClass

접근설정	정책
ReadWriteOnce	하나의 싱글노드에서만 읽고 쓰기
ReadOnlyMany	하나의 싱글노드에서 쓰고 여러 노드에게 읽기만 가능
ReadWriteMany	여러 노드에게 읽고 쓰기 가능

Kubernetes controller: Deployment

- managing pod, replica set
- Support Rolling update
- Support Rollback

Kubernetes controller: Deployment



Kubernetes controller: Deployment

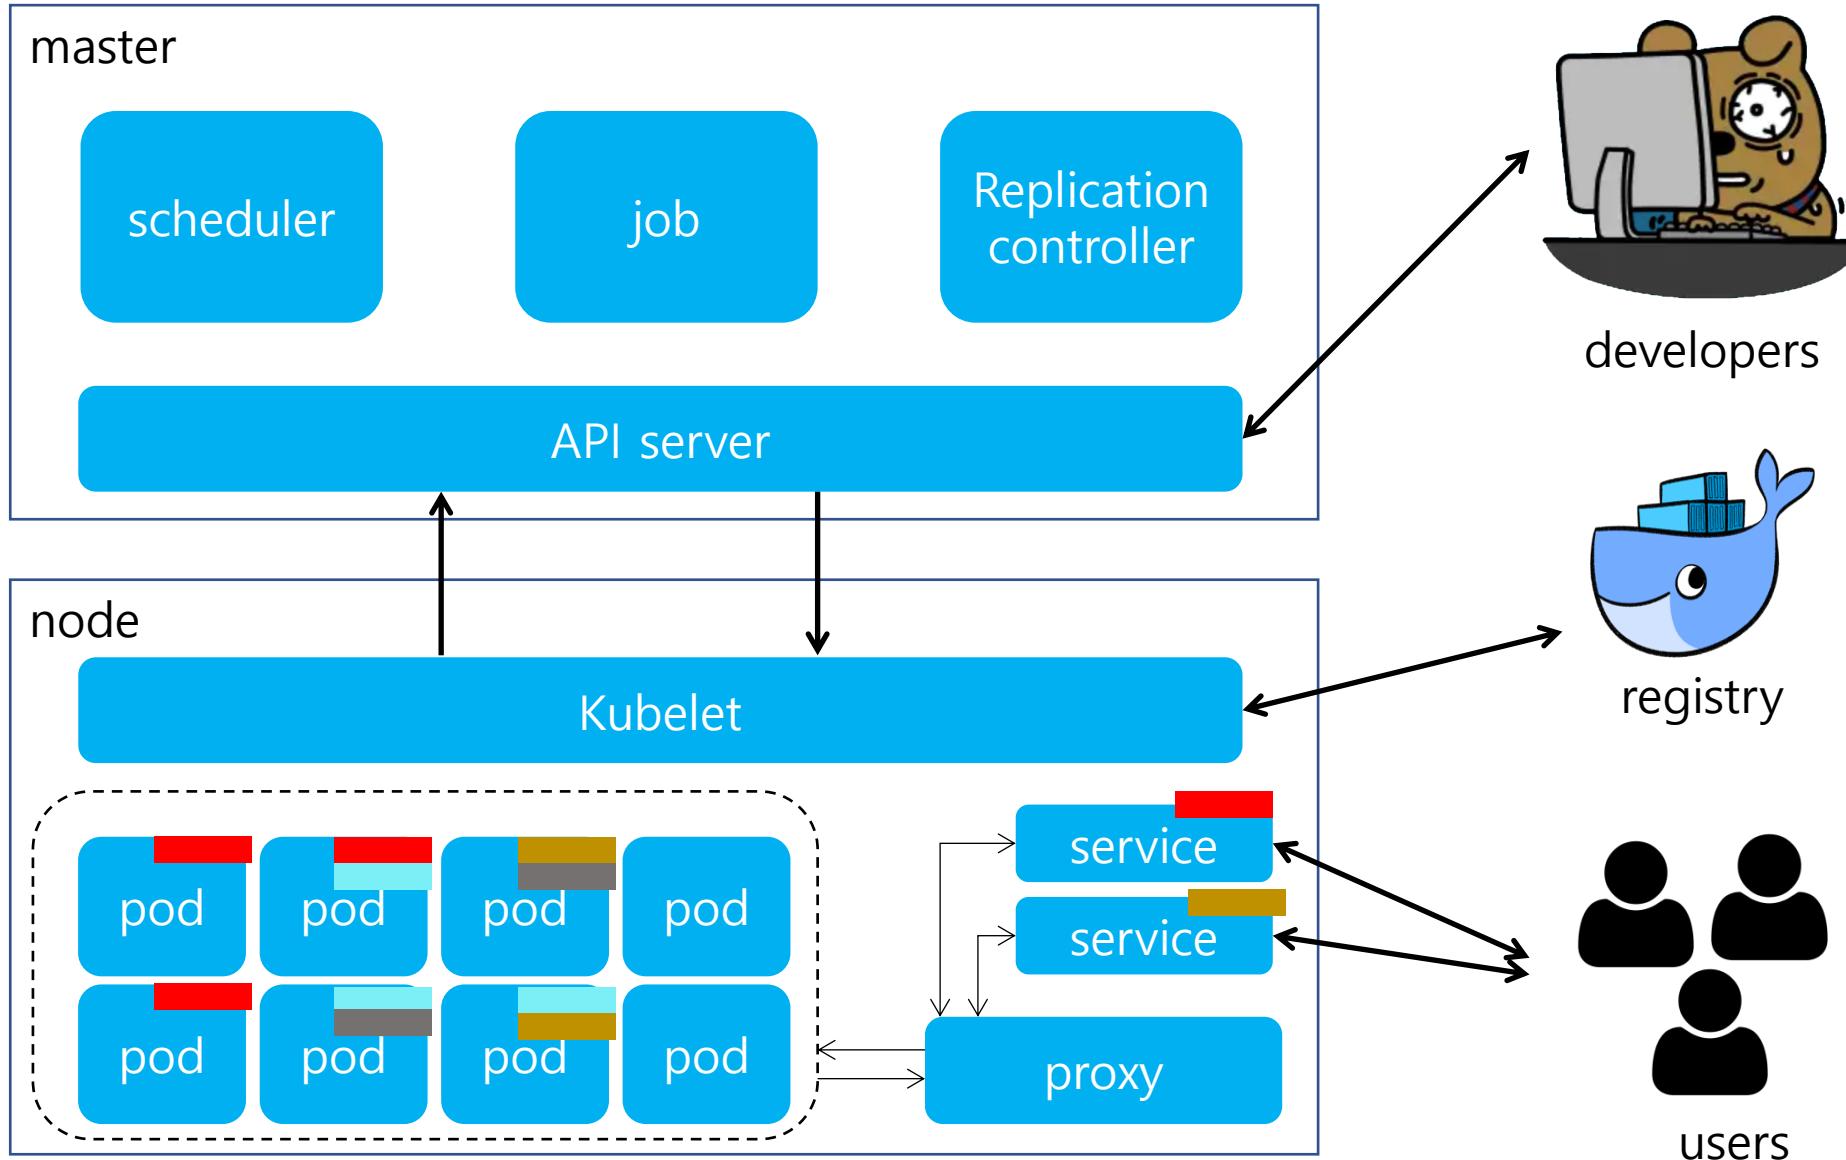
```
kind: Deployment
apiVersion: extensions/v1beta1
metadata:
  name: hostname-101-deployment
spec:
  replicas: 3
  selector:
    # Like saying "Make sure there are three pods running
    # with the label app = hostname and version = v101"
    matchLabels:
      app: hostname
      version: v101
  template:
    metadata:
      labels:
        # The `app` label is used by both the service
        # and the deployment to select the pods they operate on.
        app: hostname
        # The `version` label is used only by the deployment
        # to control replication.
        version: v101
    spec:
      containers:
        - name: nginx-hostname
          image: kubegoldenguide/nginx-hostname:1.0.1
          ports:
            - containerPort: 80
```

<https://matthewpalmer.net/kubernetes-app-developer/articles/kubernetes-deployment-tutorial-example-yaml.html>

Kubernetes Object: service

- Pod의 IP는 클러스터 내부 IP -> 그럼 외부에서는 어떻게 접속?
- service를 이용해서 클러스터 외부로 노출
- OSI7계층 중 3계층(네트워크계층 : ip부여 / 라우팅)
- 서비스에 속한 Pod는 label selector로 판단
- 로드밸런싱 및 서비스 디스커버리 수행

Kubernetes Object: service

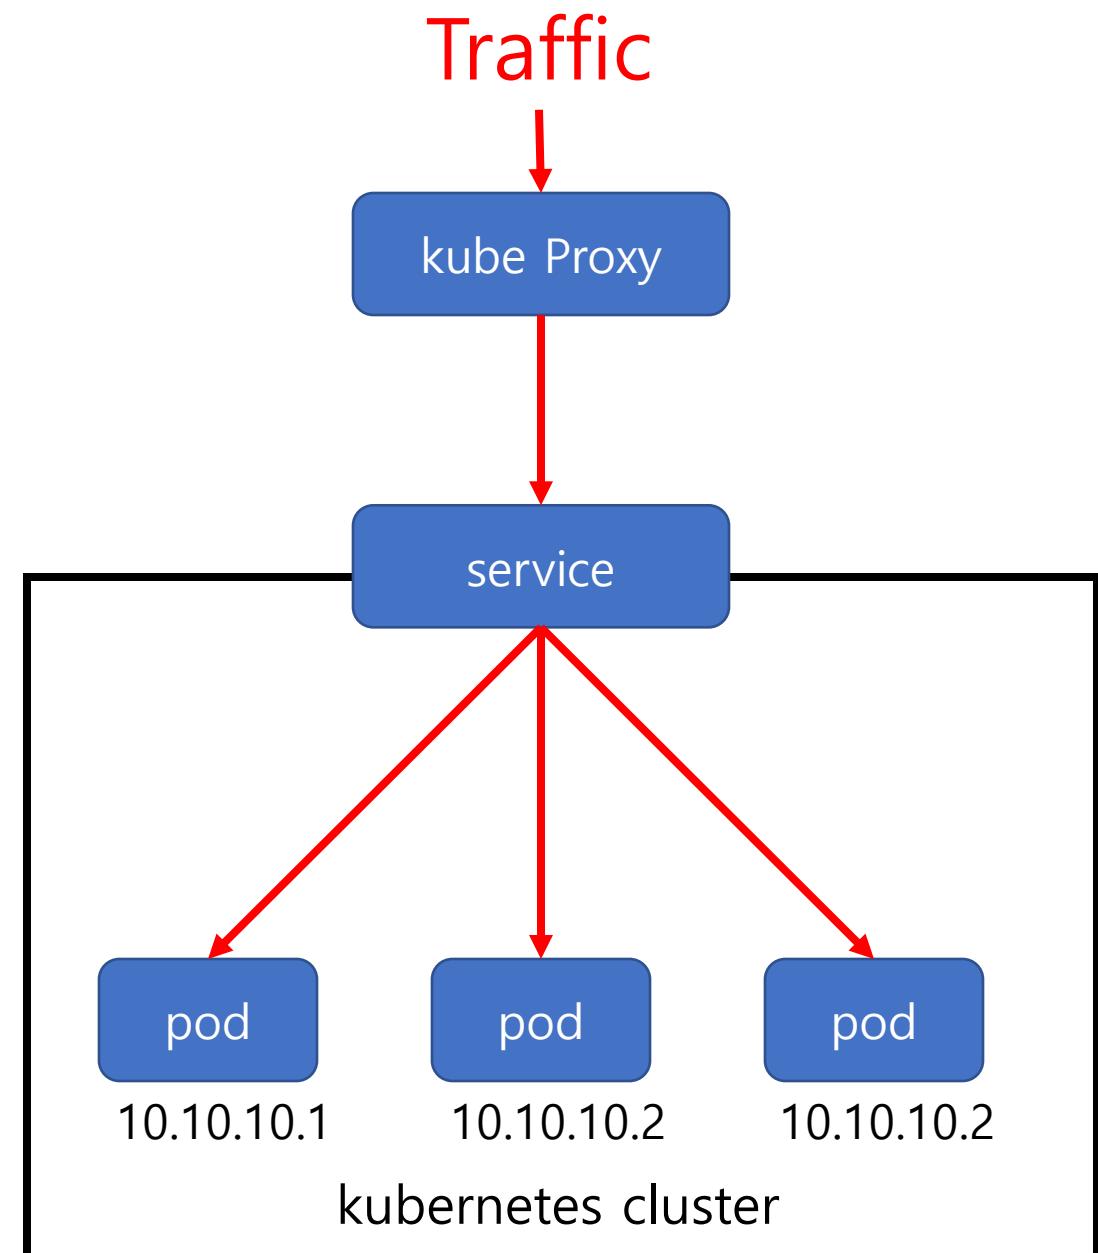


Cluster IP

- Cluster 내부 통신을 위한 IP
- 외부에서 Access 불가능
- 외부에서 접근을 위해 kube proxy 이용

```
apiVersion: v1
kind: Service
metadata:
  name: my-internal-service
spec:
  selector:
    app: my-app
  type: ClusterIP
  ports:
  - name: http
    port: 80
    targetPort: 80
    protocol: TCP
```

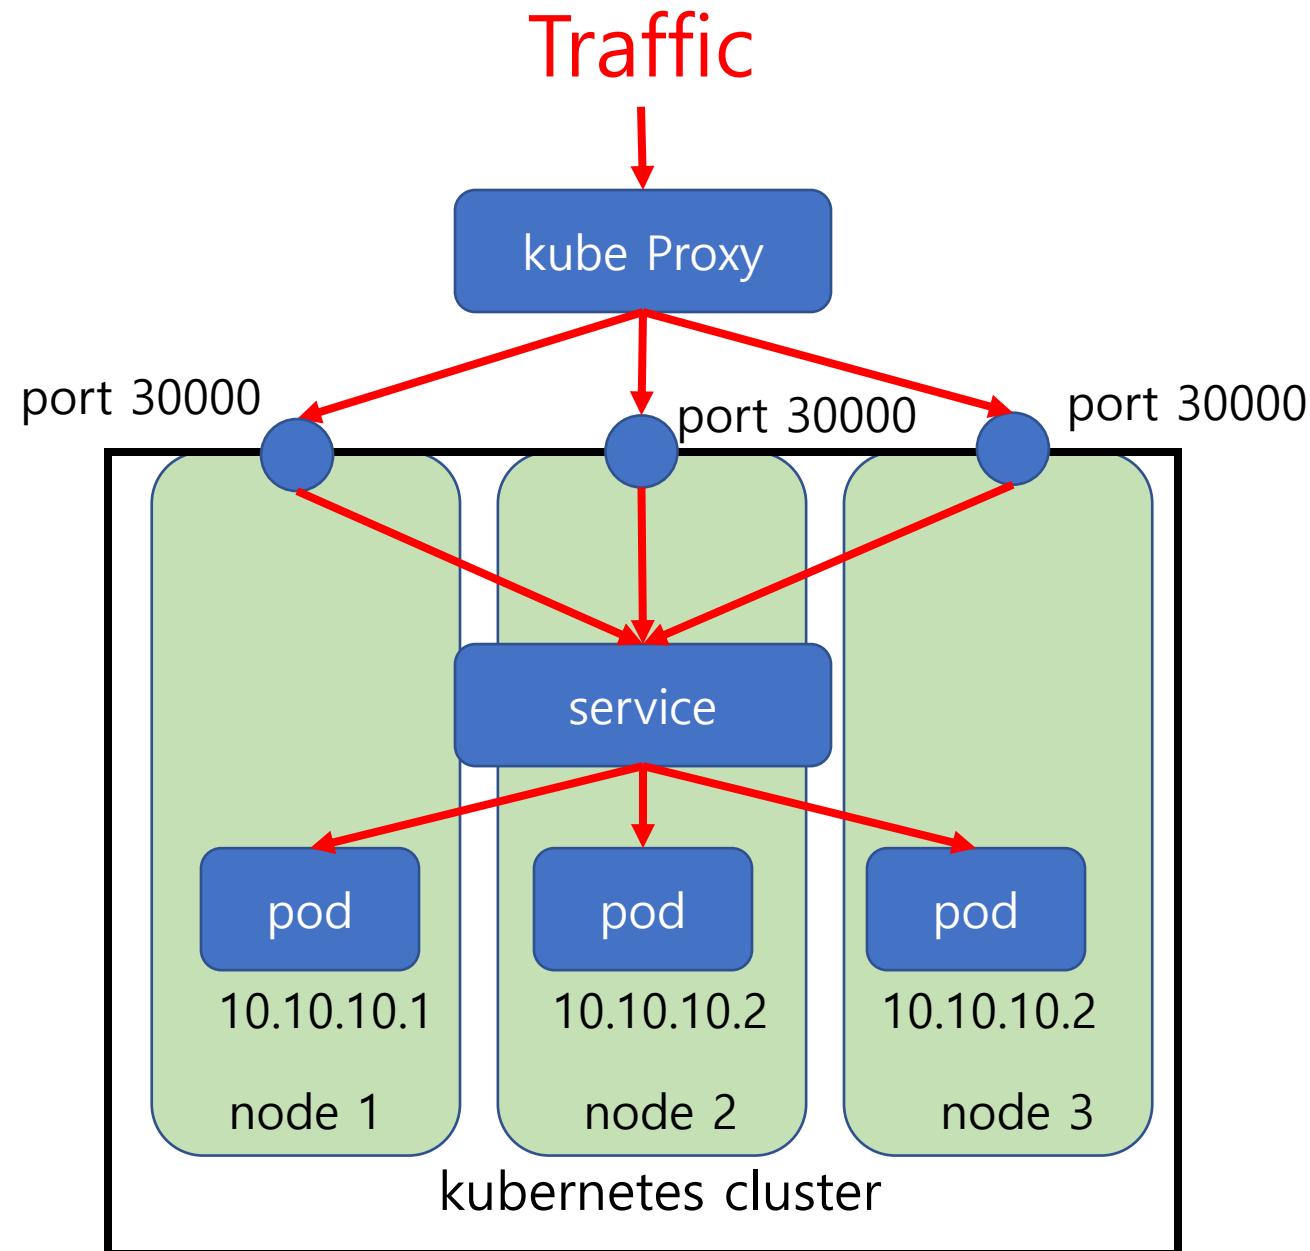
```
$ kubectl proxy --port=8080
```



Node Port

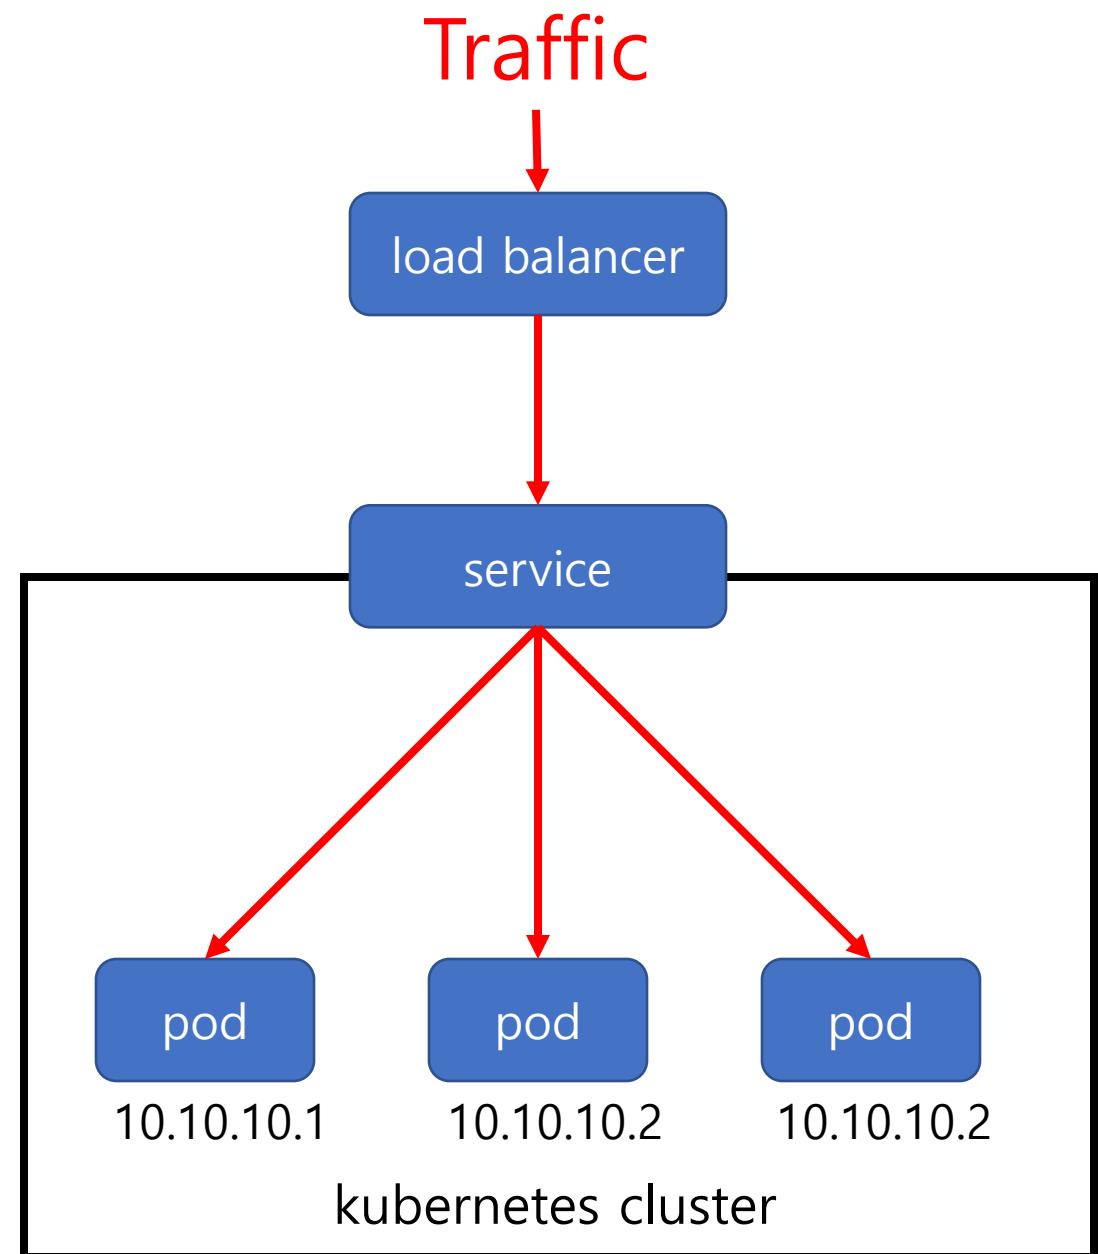
- NAT을 사용하여 각 노드의 Port 오픈
- 30000 ~ 32767 범위의 포트 할당
- 명시적 지정이 아니면 랜덤 할당

```
apiVersion: v1
kind: Service
metadata:
  name: my-nodeport-service
spec:
  selector:
    app: my-app
  type: NodePort
  ports:
  - name: http
    port: 80
    targetPort: 80
    nodePort: 30036
    protocol: TCP
```



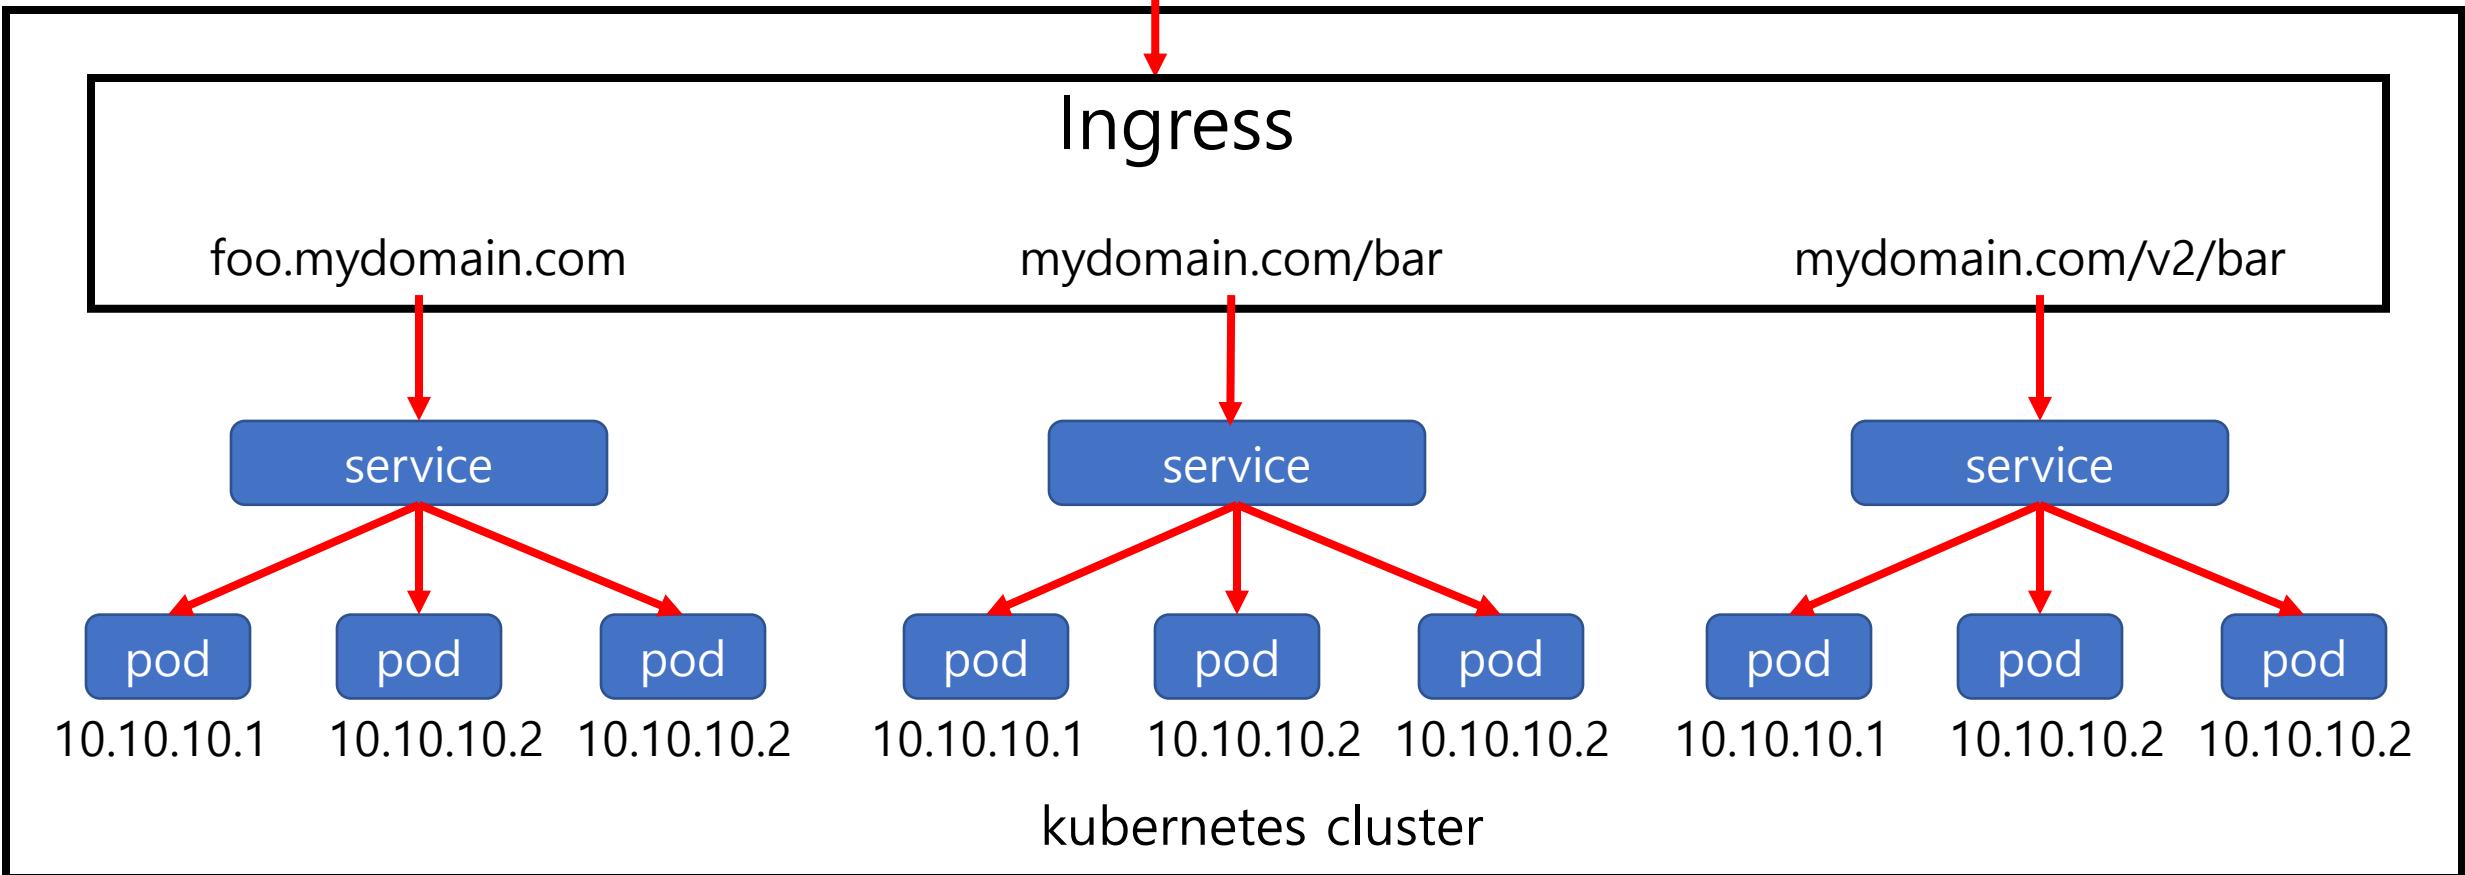
Load Balancer

- External 로드밸런서를 통해 고정된 외부 IP 배정



Ingress

Traffic



라벨 (label)

- key-value 쌍의 객체 집합이며 객체의 식별에 사용
- Pod를 그룹화할 때 사용
- 복제 컨트롤러(Replication Controller), 복제 집합(Replica set) 등에서 사용

라벨 셀렉터

- 라벨을 기반으로 객체를 선택
- 동일한 라벨이 부여된 객체들을 연산자를 통해 일괄 선택 가능
 - ex) 라벨 키가 role이고 값이 webserver인 객체 선택
 - role = webserver
 - = 또는 ==, != 을 통해 라벨에 따른 객체 선택
 - 복수의 조건은 쉼표로 구분
 - role = webserver, application != foo
 - 집합은 다양한 값을 기반으로 객체 선택
 - role in (webserver, backend)